
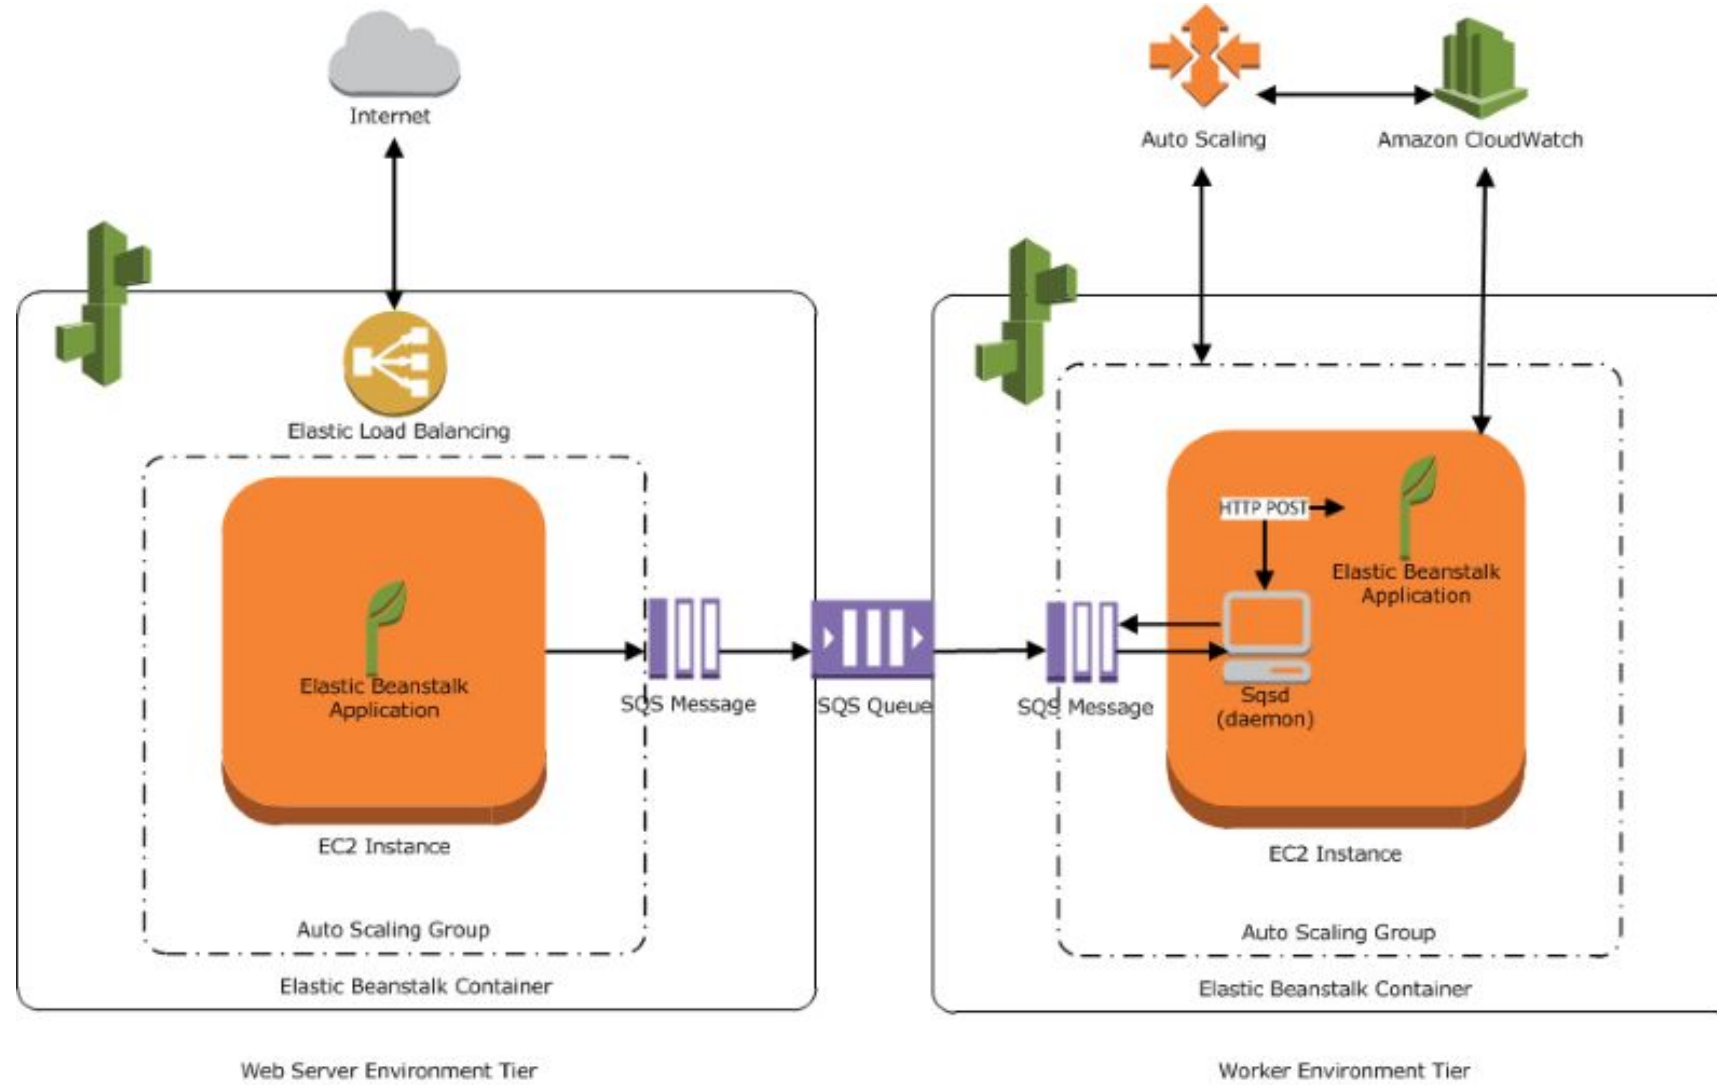




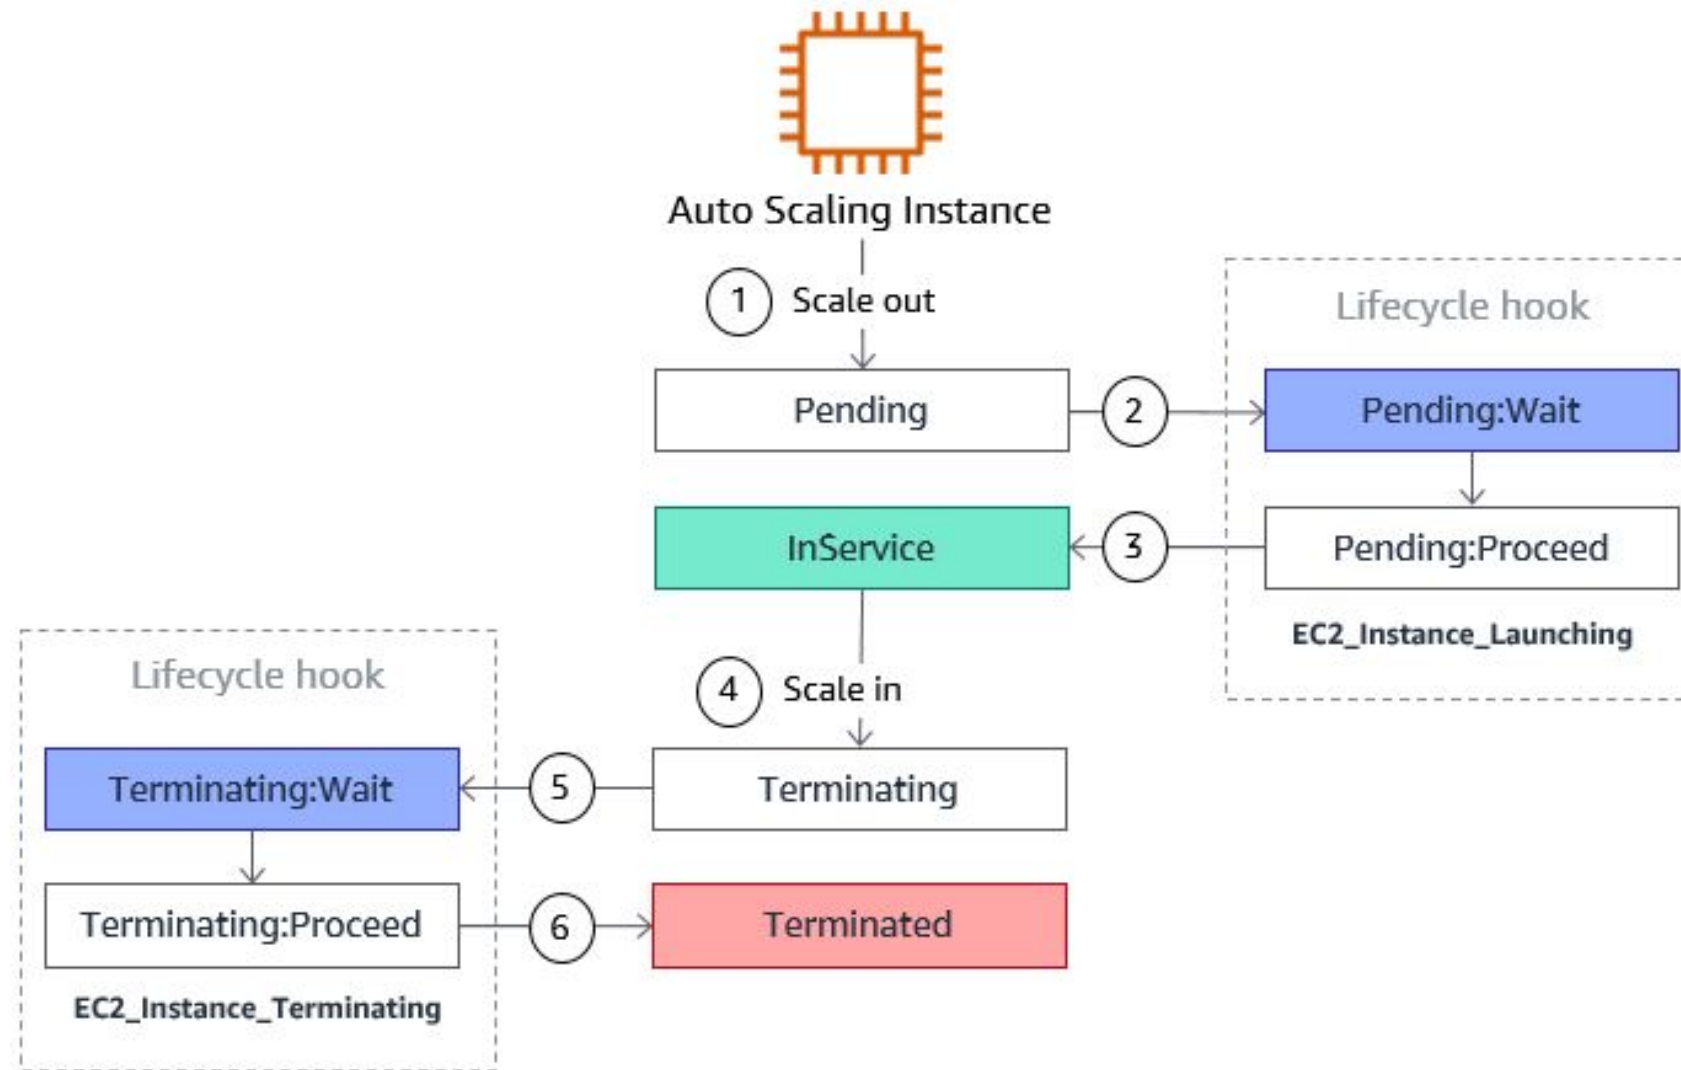
# CSAA Test 1







<https://docs.aws.amazon.com/autoscaling/ec2/userguide/lifecycle-hooks.html>



When a scale-in event occurs, a lifecycle hook pauses the instance before it is terminated and sends you a notification using Amazon EventBridge. While the instance is in the wait state, you can invoke an AWS Lambda function or connect to the instance to download logs or other data before the instance is fully terminated.



**AZ-1**

**AZ-2**

**AZ-3**

<b>A</b>	<b>2</b>	<b>2</b>	<b>2</b>
<b>B</b>	<b>3</b>	<b>3</b>	<b>0</b>
<b>C</b>	<b>4</b>	<b>2</b>	<b>2</b>
<b>D</b>	<b>6</b>	<b>6</b>	<b>0</b>
<b>E</b>	<b>3</b>	<b>3</b>	<b>3</b>



# LAMBDA



# S3 EVENT NOTIFICATION

## Events

+ Add notification

Delete

Edit

Name	Events	Filter	Type
New event			

Name

e.g. MyEmailEventForPut

Events

☐ RRSObjectLost

☐ Put

☐ Post

☐ Copy

☐ Complete Multipart Upload

☐ Delete

☐ Delete Marker Created

☐ ObjectCreate (All)

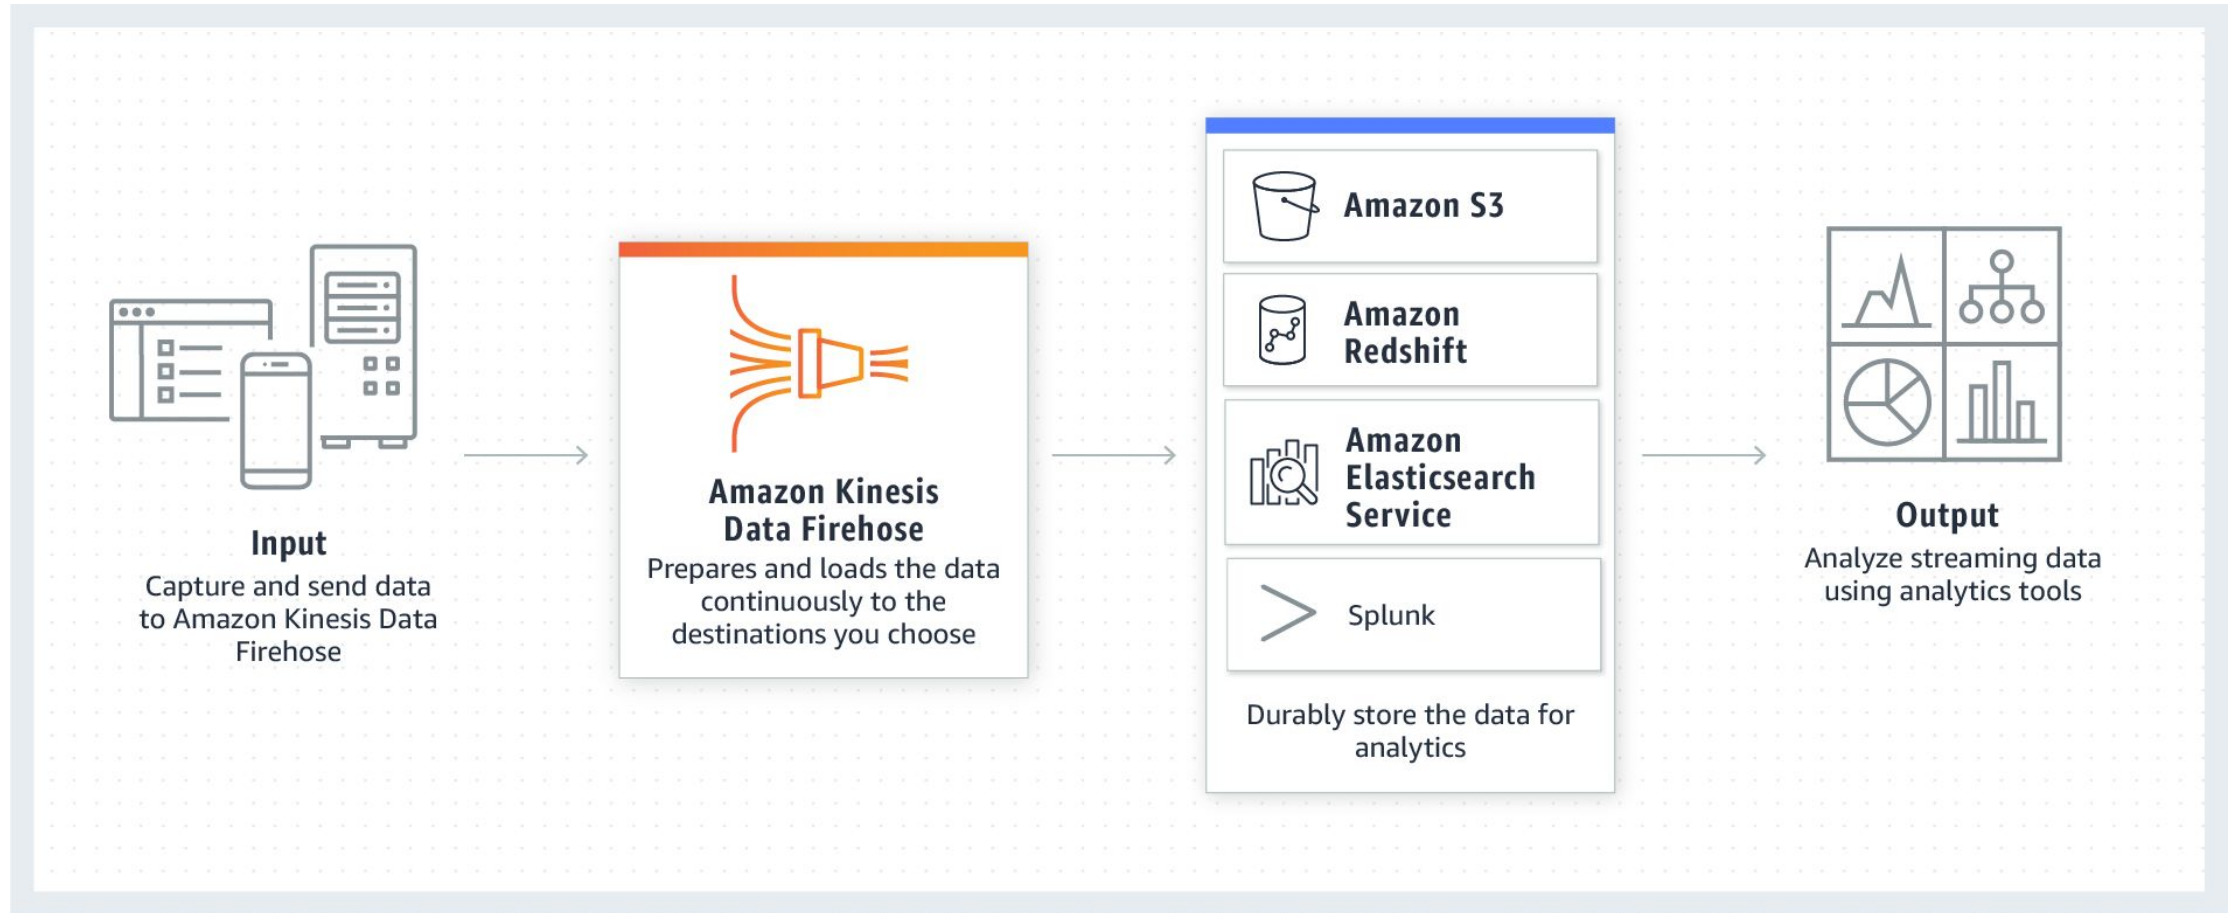
☐ ObjectDelete (All)

CLARUSWAY  
WAY TO REINVENT YOURSELF

6



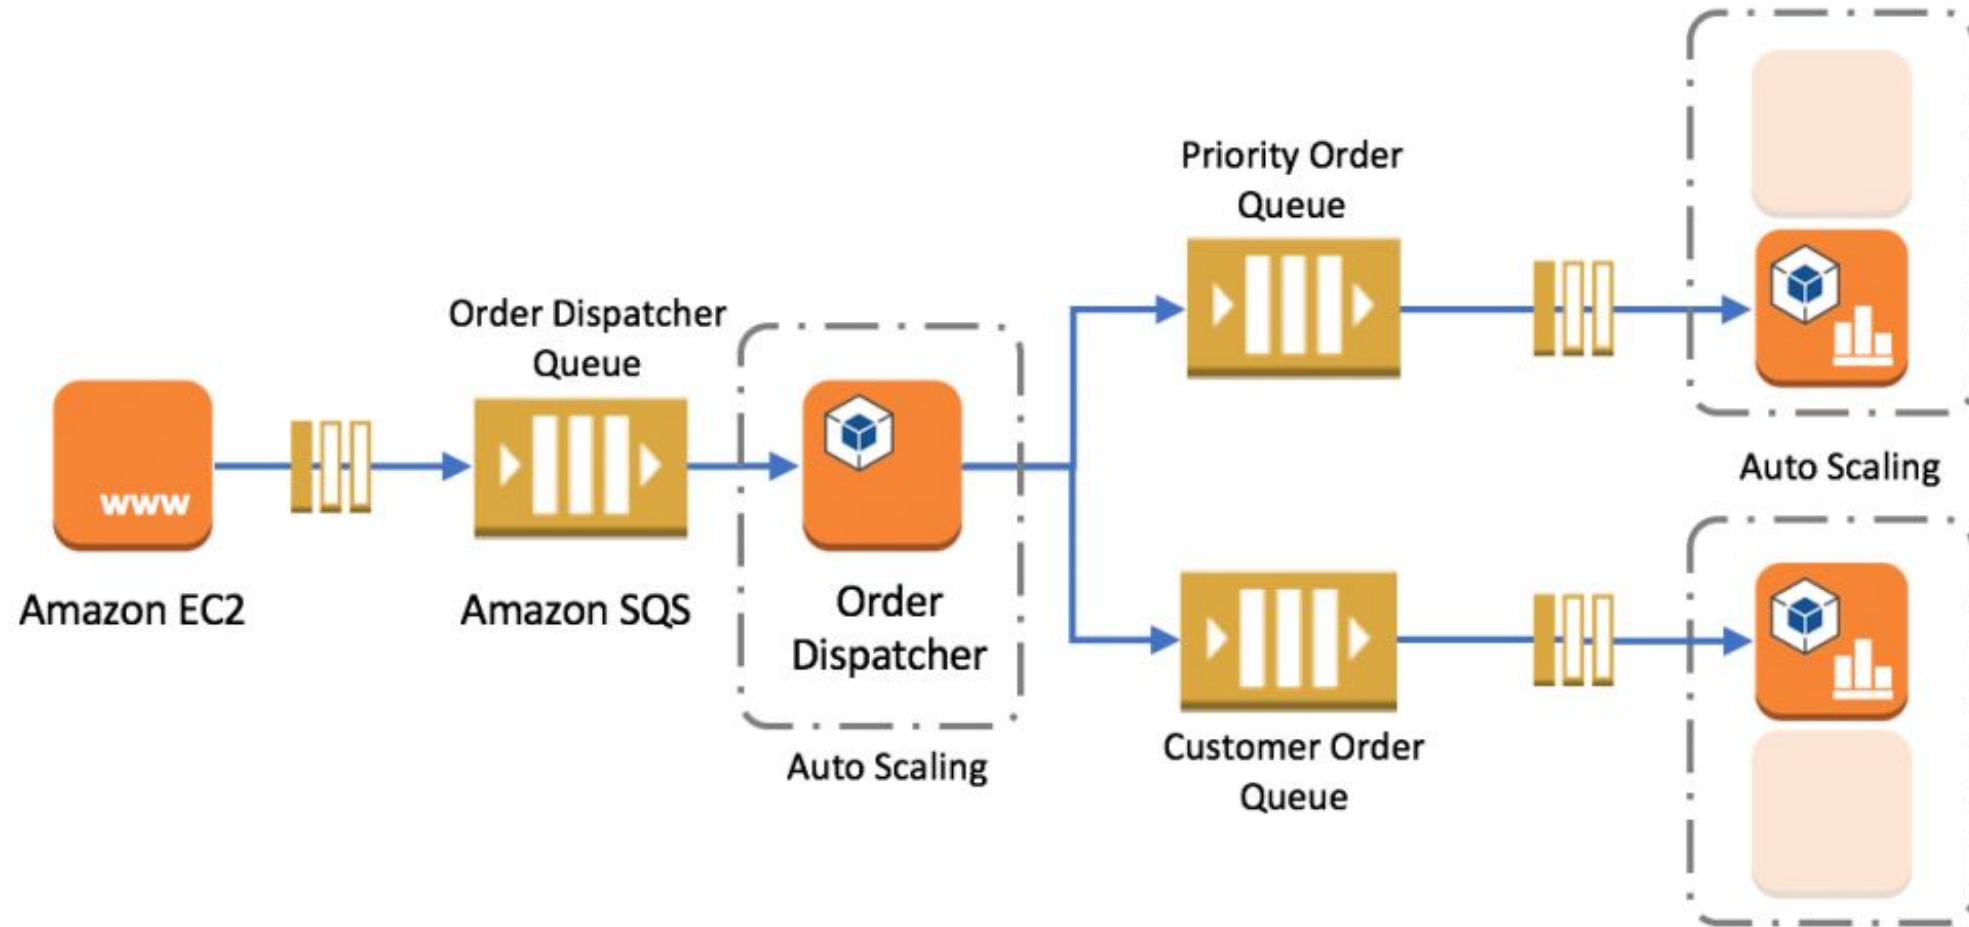
# KINESIS



**KINESIS is used to collect, process and analyze real-time data**

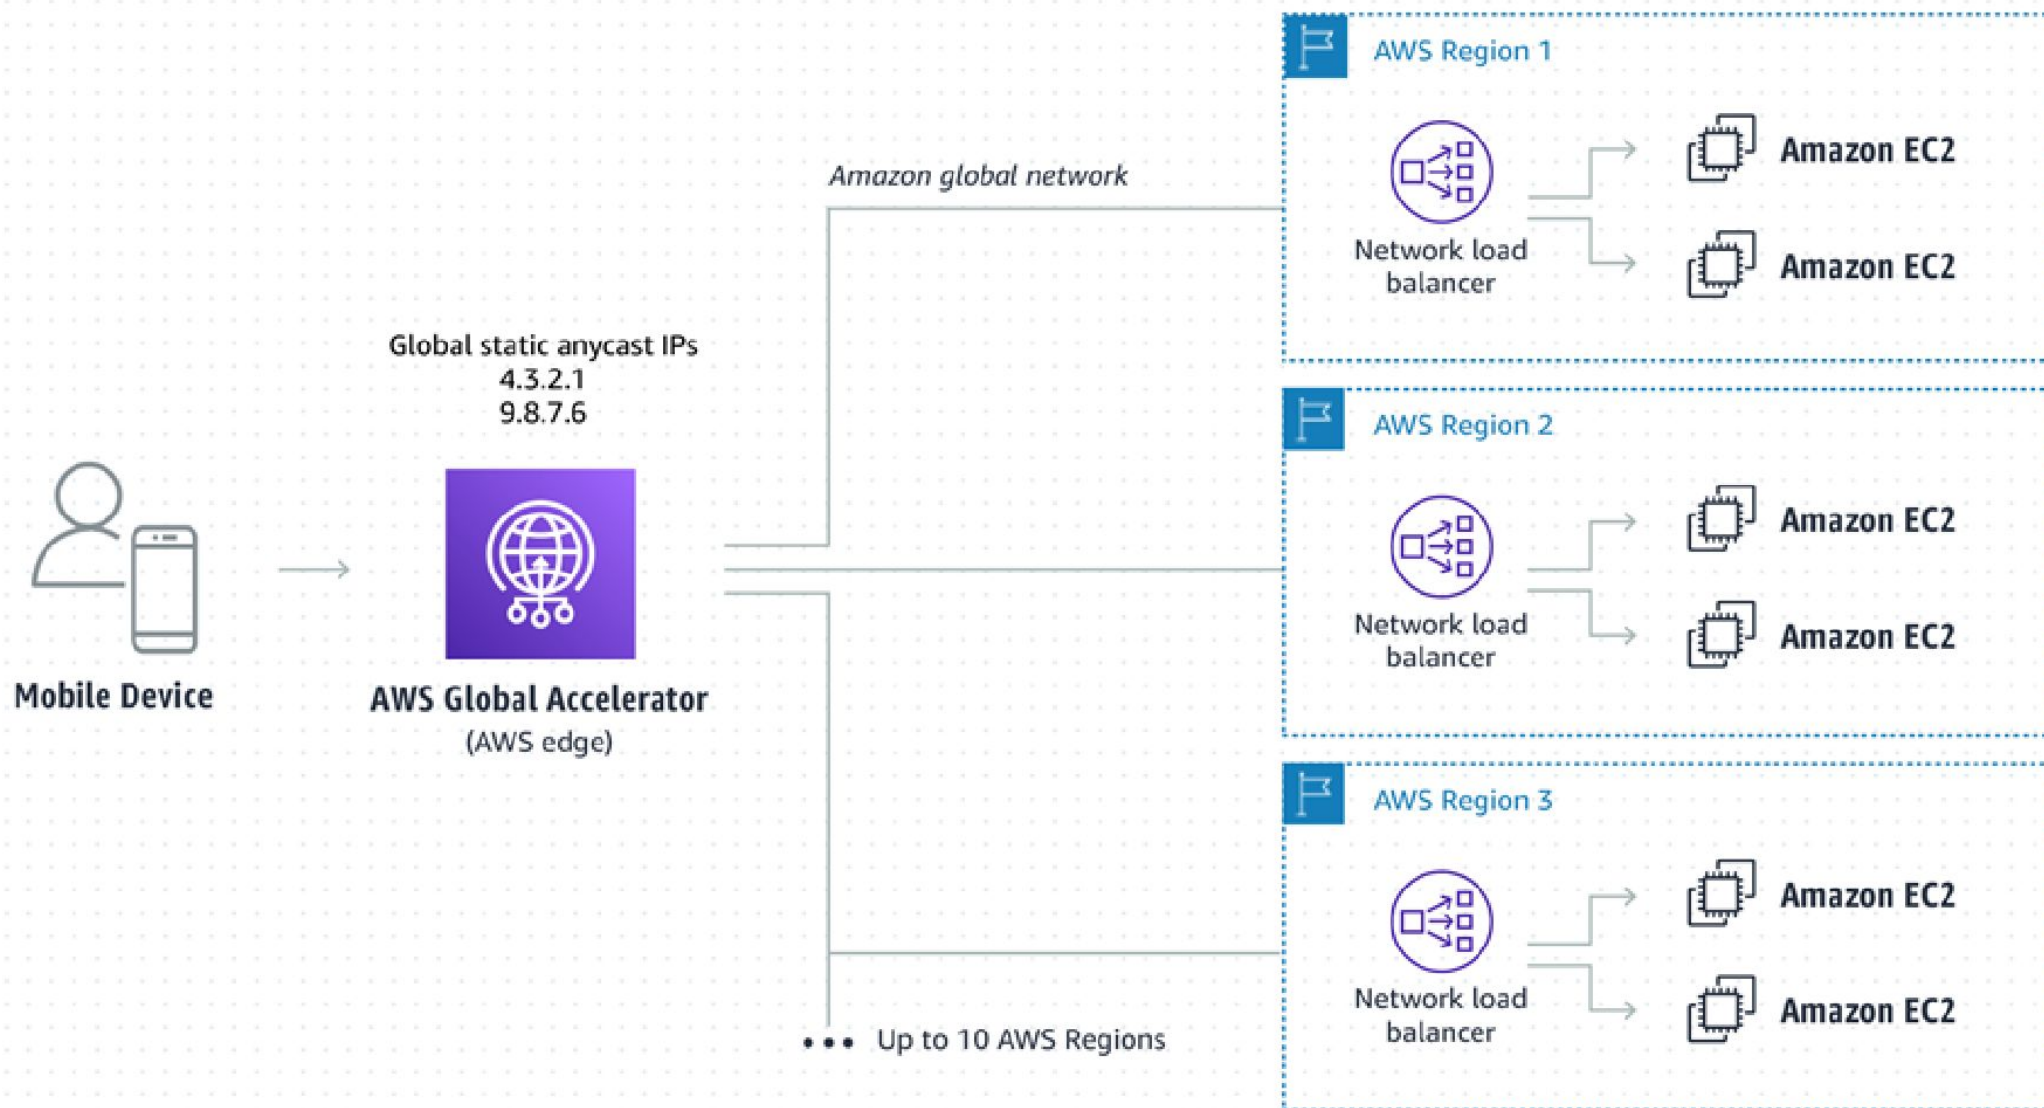


# SQS



**SQS MAX FILE SIZE = 2 GB**







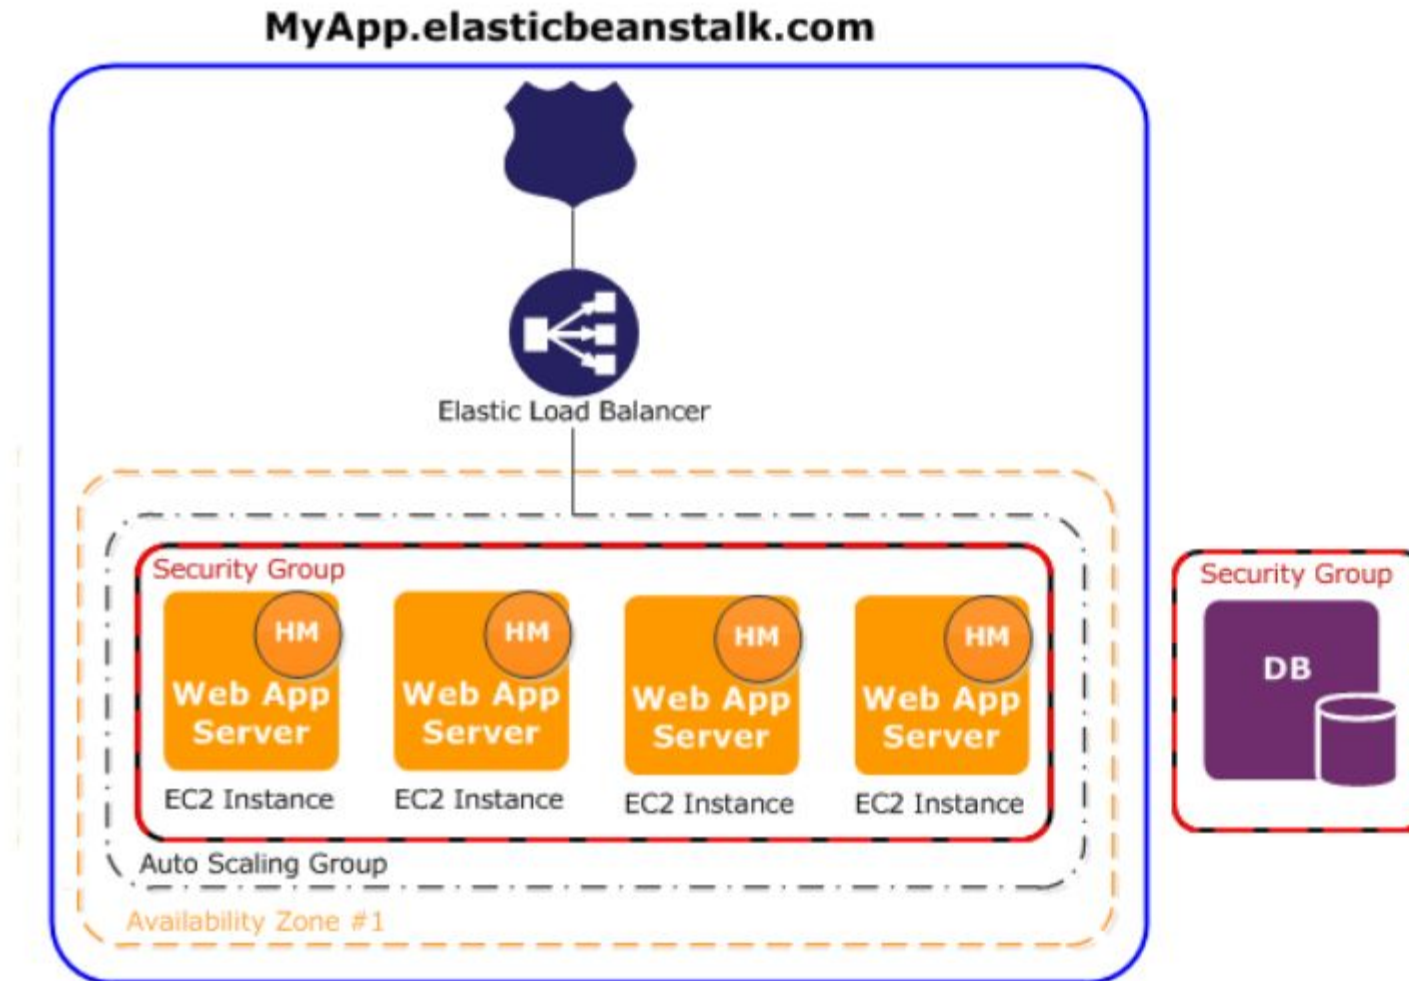
**Q: How is AWS Global Accelerator different from a DNS-based traffic management solution?**

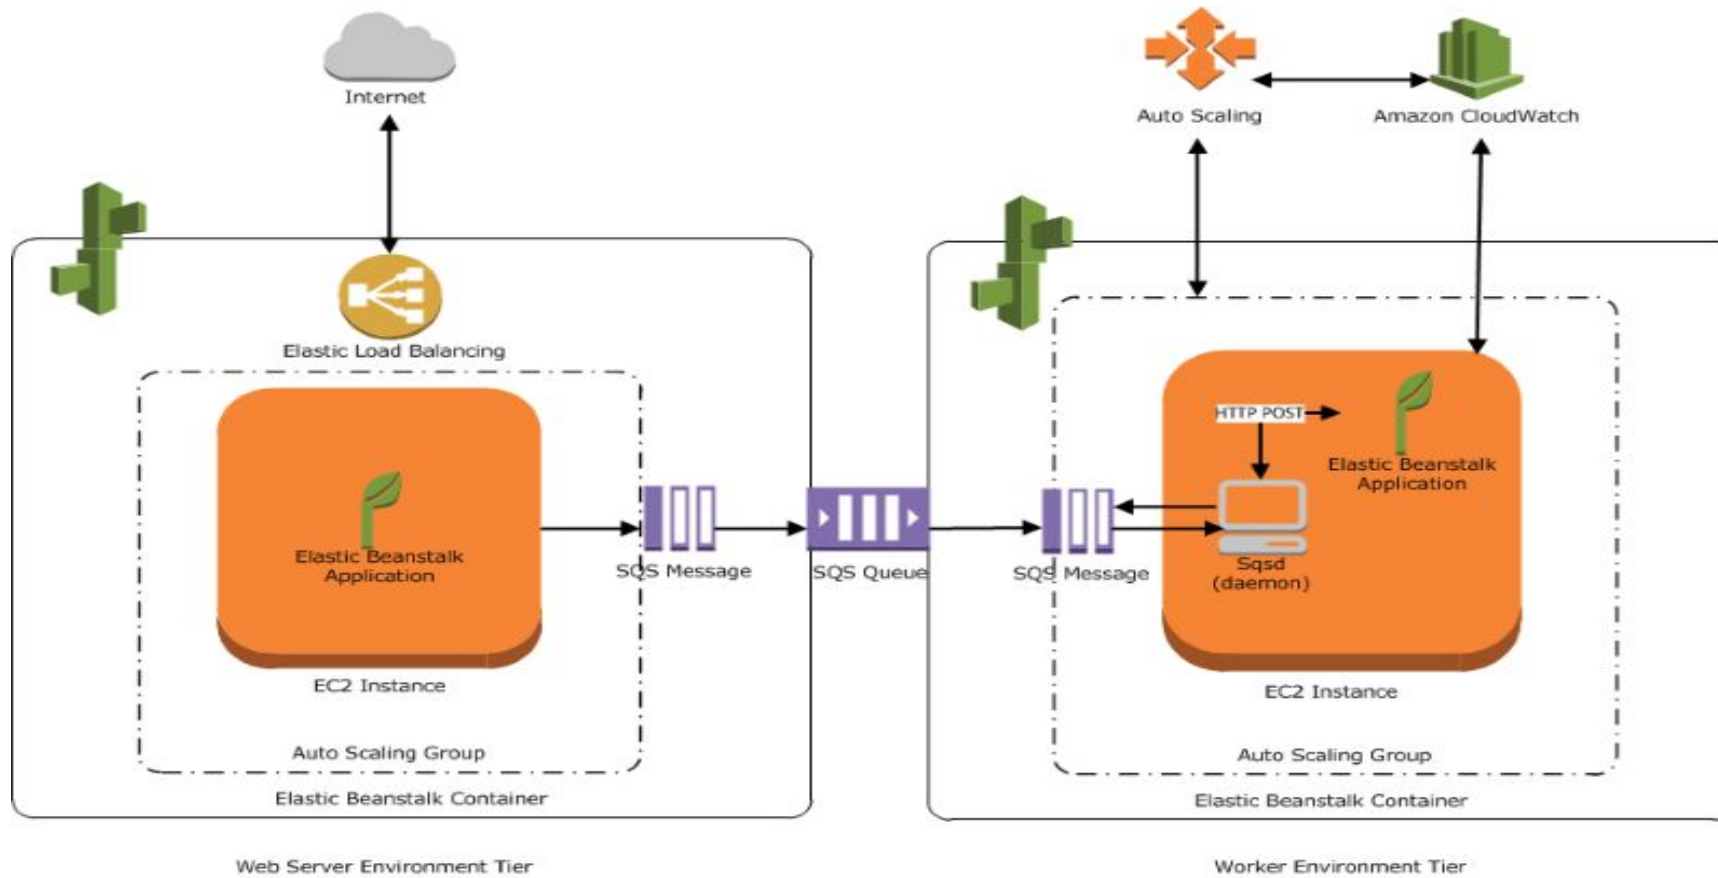
A: First, some client devices and internet resolvers cache DNS answers for long periods of time. So when you make a configuration update, or there's an application failure or change in your routing preference, you don't know how long it will take before all of your users receive updated IP addresses. With AWS Global Accelerator, you don't have to rely on the IP address caching settings of client devices. Change propagation takes a matter of seconds, which reduces your application downtime. Second, with Global Accelerator, you get static IP addresses that provide a fixed entry point to your applications. This lets you easily move your endpoints between Availability Zones or between AWS Regions, without having to update the DNS configuration or client-facing applications.

---



This following diagram illustrates an example Elastic Beanstalk architecture for a web server environment tier and shows how the components in that type of environment tier work together. The remainder of this section discusses all the components in more detail.





## Supported platform versions

- Docker
- Multicontainer Docker
- Preconfigured Docker
- Go
- Java SE
- Tomcat
- .NET Core on Linux
- .NET on Windows Server
- Node.js
- PHP
- Python
- Ruby



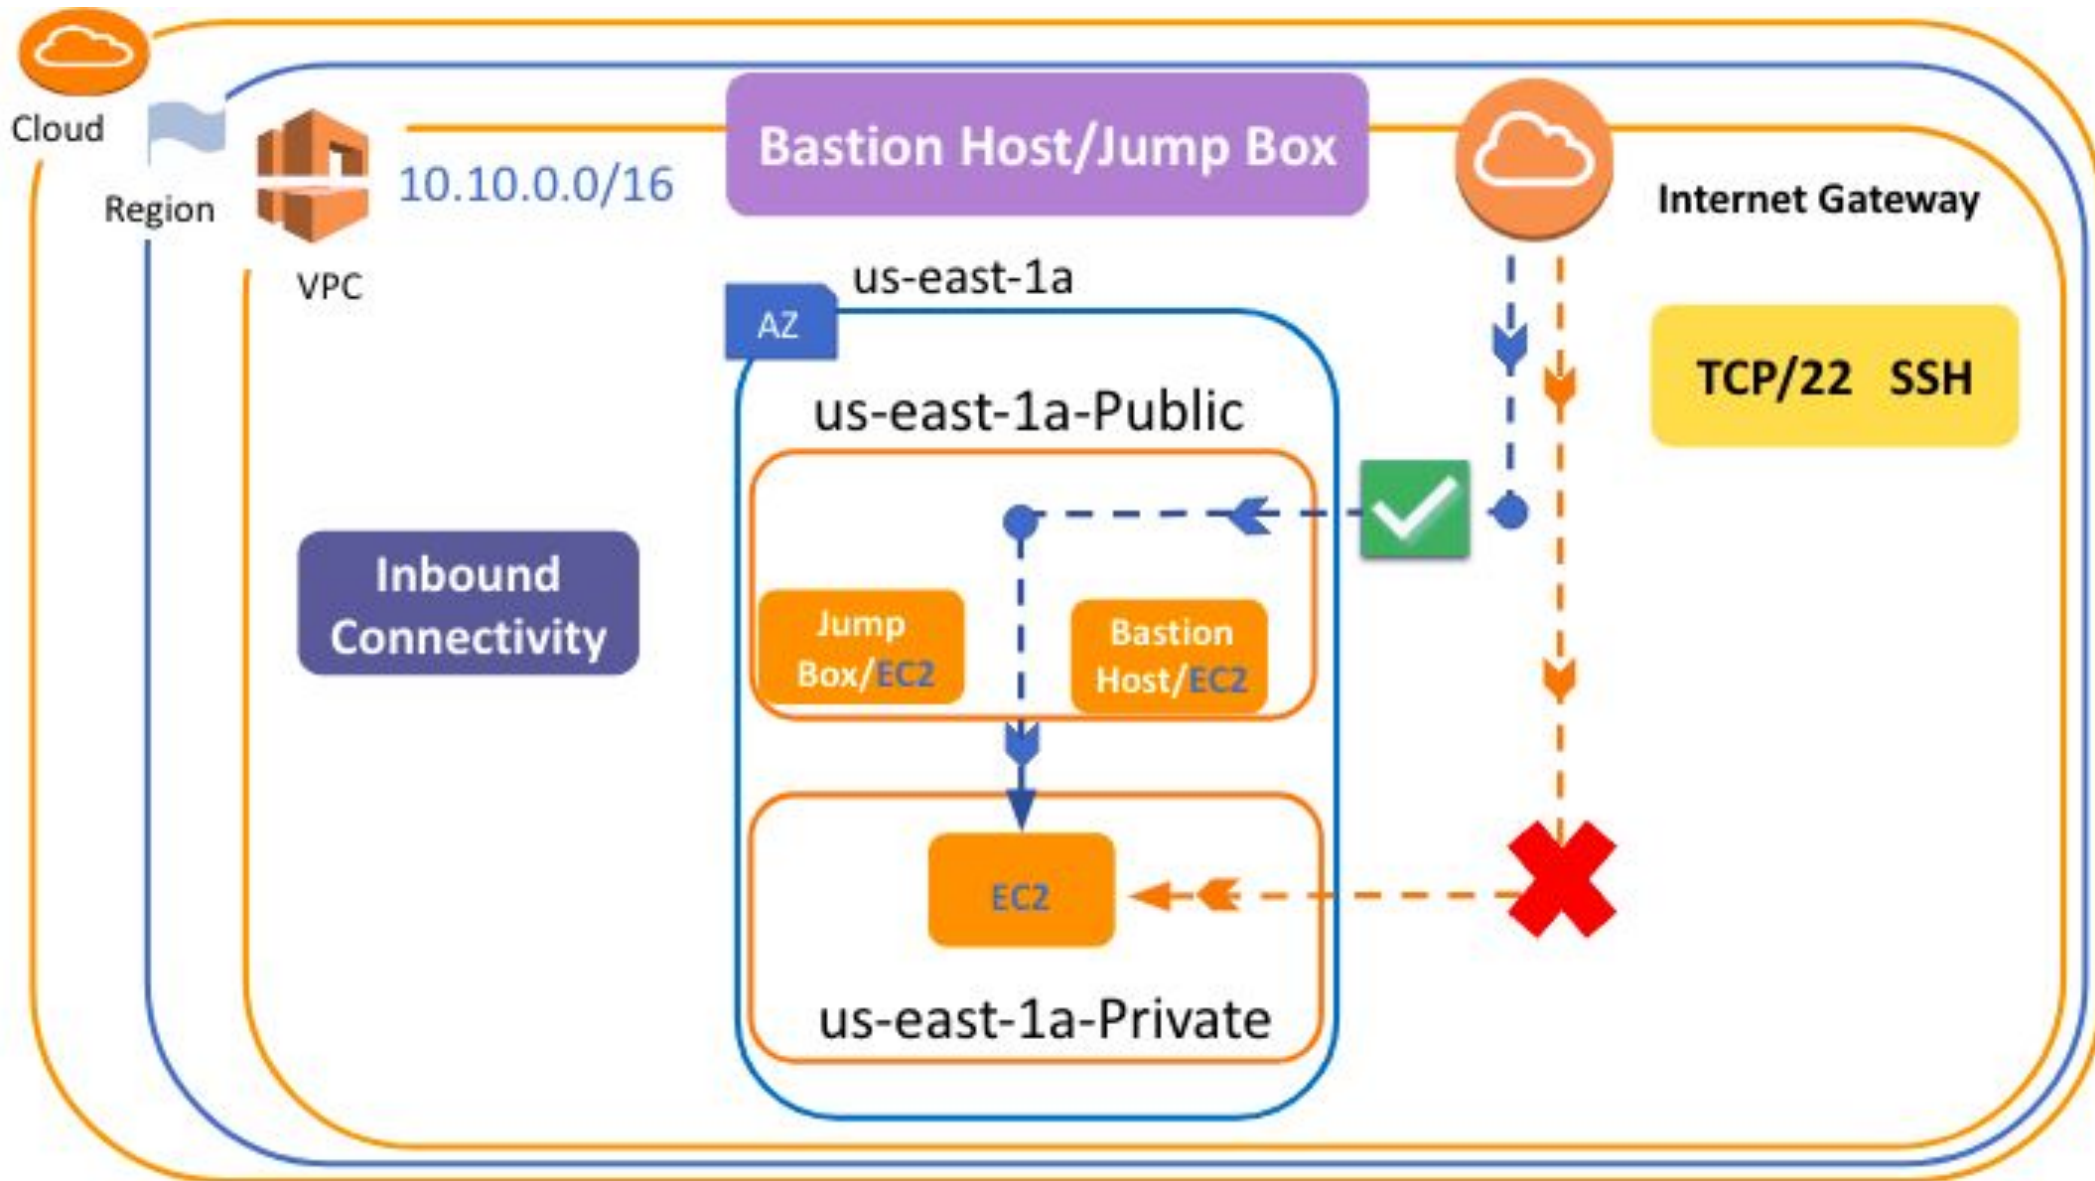


	Throughput Optimized HDD	Cold HDD
Volume type	st1	sc1
Durability	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)
Use cases	<ul style="list-style-type: none"><li>• Big data</li><li>• Data warehouses</li><li>• Log processing</li></ul>	<ul style="list-style-type: none"><li>• Throughput-oriented storage for data that is infrequently accessed</li><li>• Scenarios where the lowest storage cost is important</li></ul>

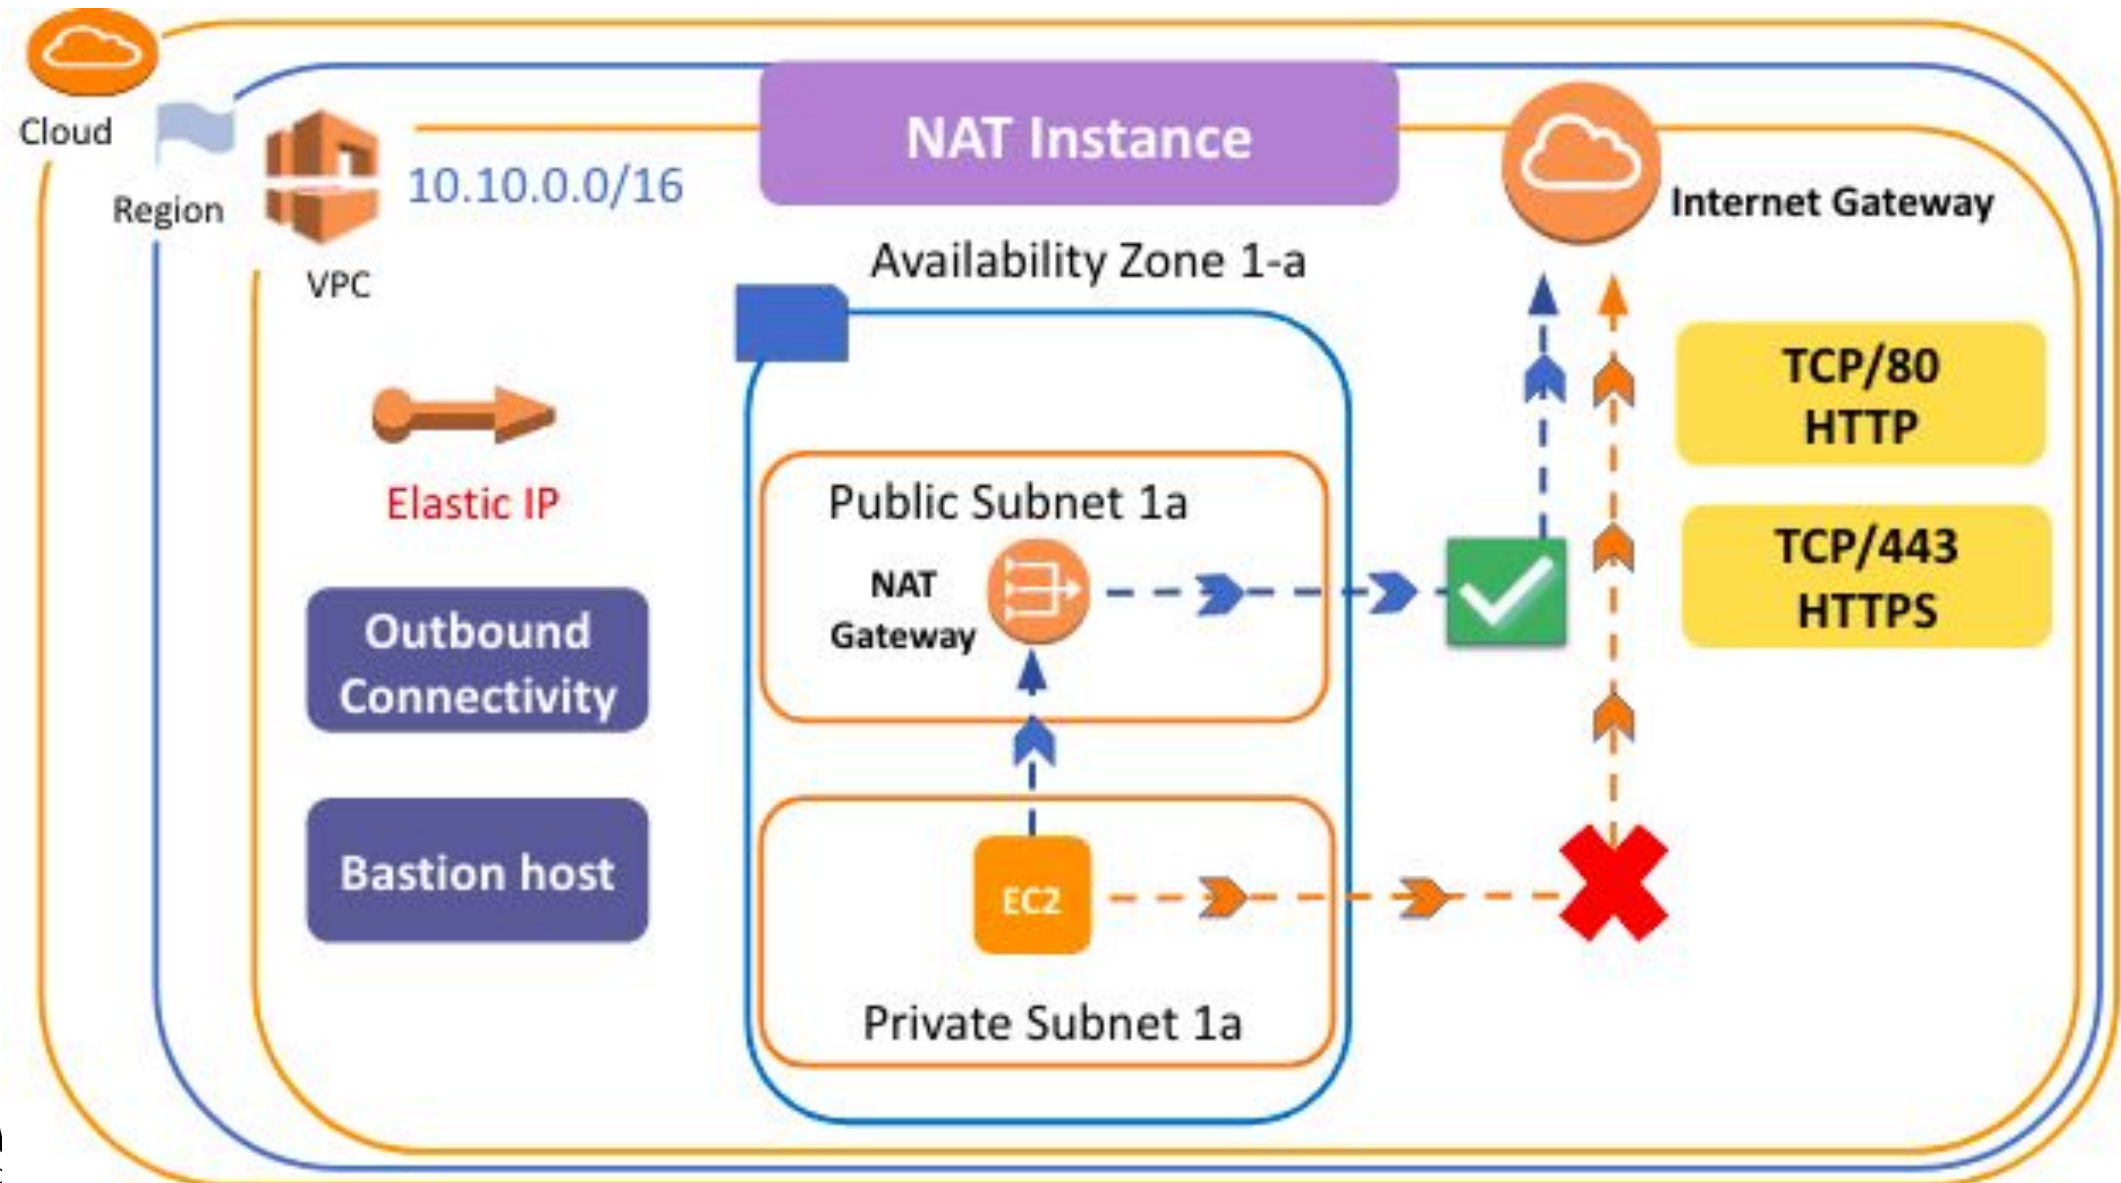
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volume-types.html>

<https://aws.amazon.com/ebs/features/>

# BASTION HOST

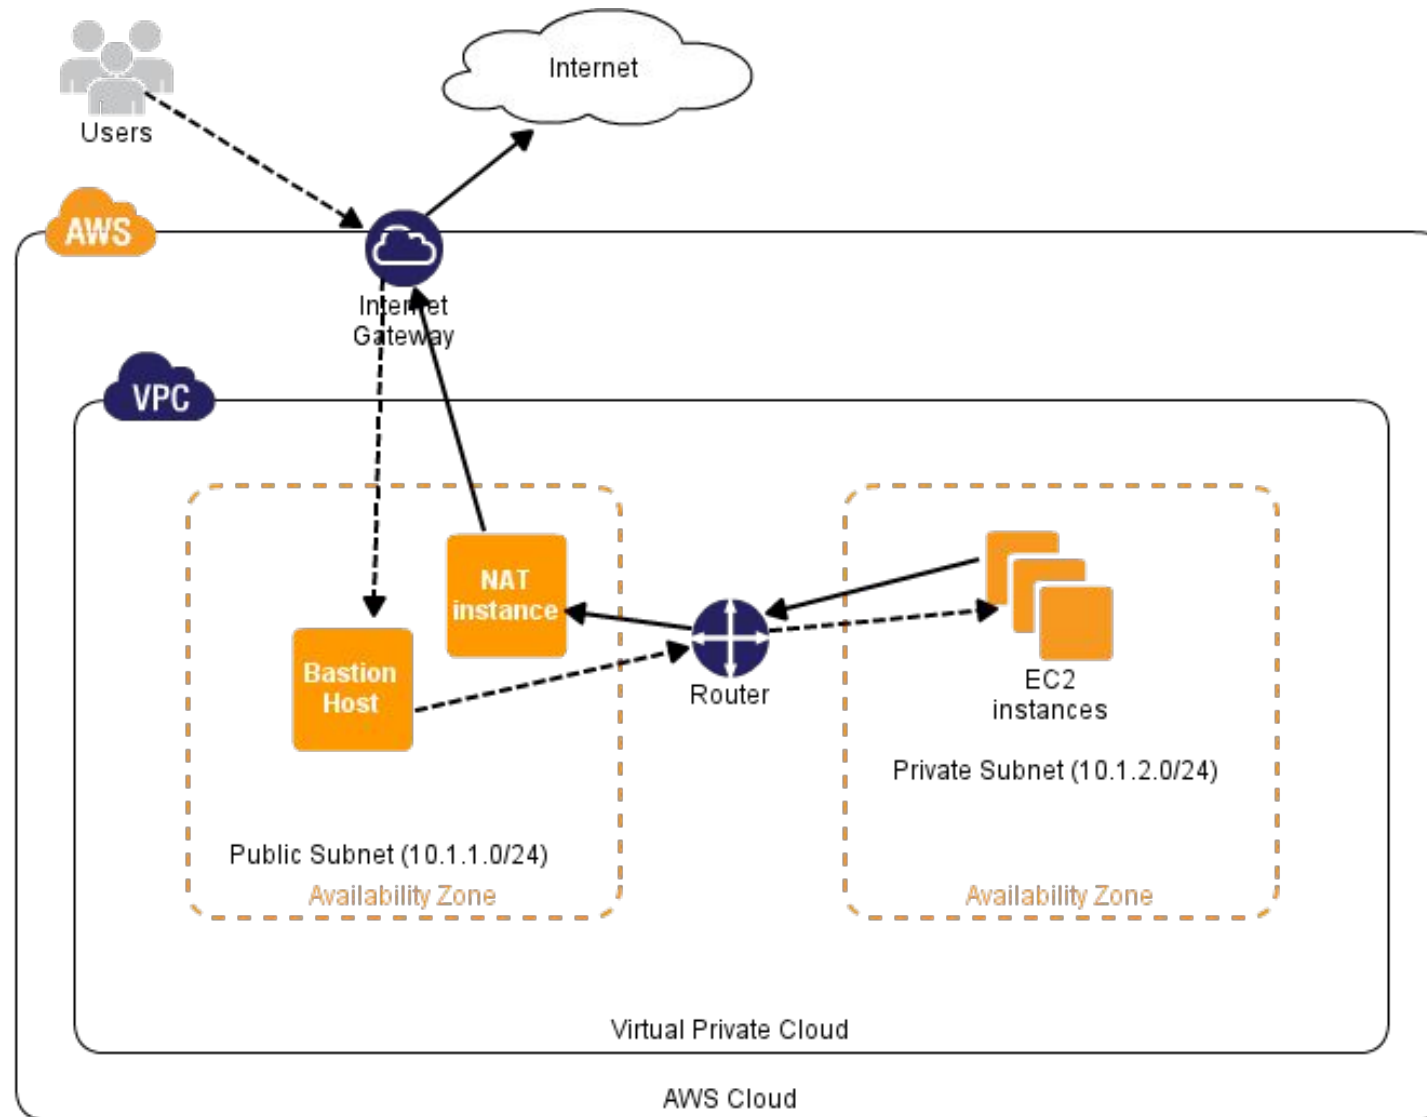


# NAT INSTANCE





# BASTION HOST VS NAT INSTANCE







# S3 Glacier Archives – Retrieval Options

<b>Expedited</b>	Expedited retrievals allow you to quickly access your data stored in the S3 Glacier Flexible Retrieval storage class or S3 Intelligent-Tiering Archive Access tier when occasional urgent requests for a subset of archives are required. For all but the largest archives (250 MB+), data accessed using <b>Expedited retrievals are typically made available within 1–5 minutes.</b>
<b>Standard</b>	Standard retrievals allow you to access any of your archives within several hours. <b>Standard retrievals typically complete within 3–5 hours.</b> This is the default option for retrieval requests that do not specify the retrieval option.
<b>Bulk</b>	Bulk retrievals are S3 Glacier’s lowest-cost retrieval option, which you can use to retrieve large amounts, even petabytes, of data inexpensively in a day. <b>Bulk retrievals typically complete within 5–12 hours.</b>

## Hybrid cloud storage challenges



### On-premises

I want to run my existing applications without change (databases, files, backups)...

I need local access to data in the cloud...

Compliance requires that I integrate with enterprise security and management tools...

I need reliable, easy connectivity to the cloud...

### In-cloud



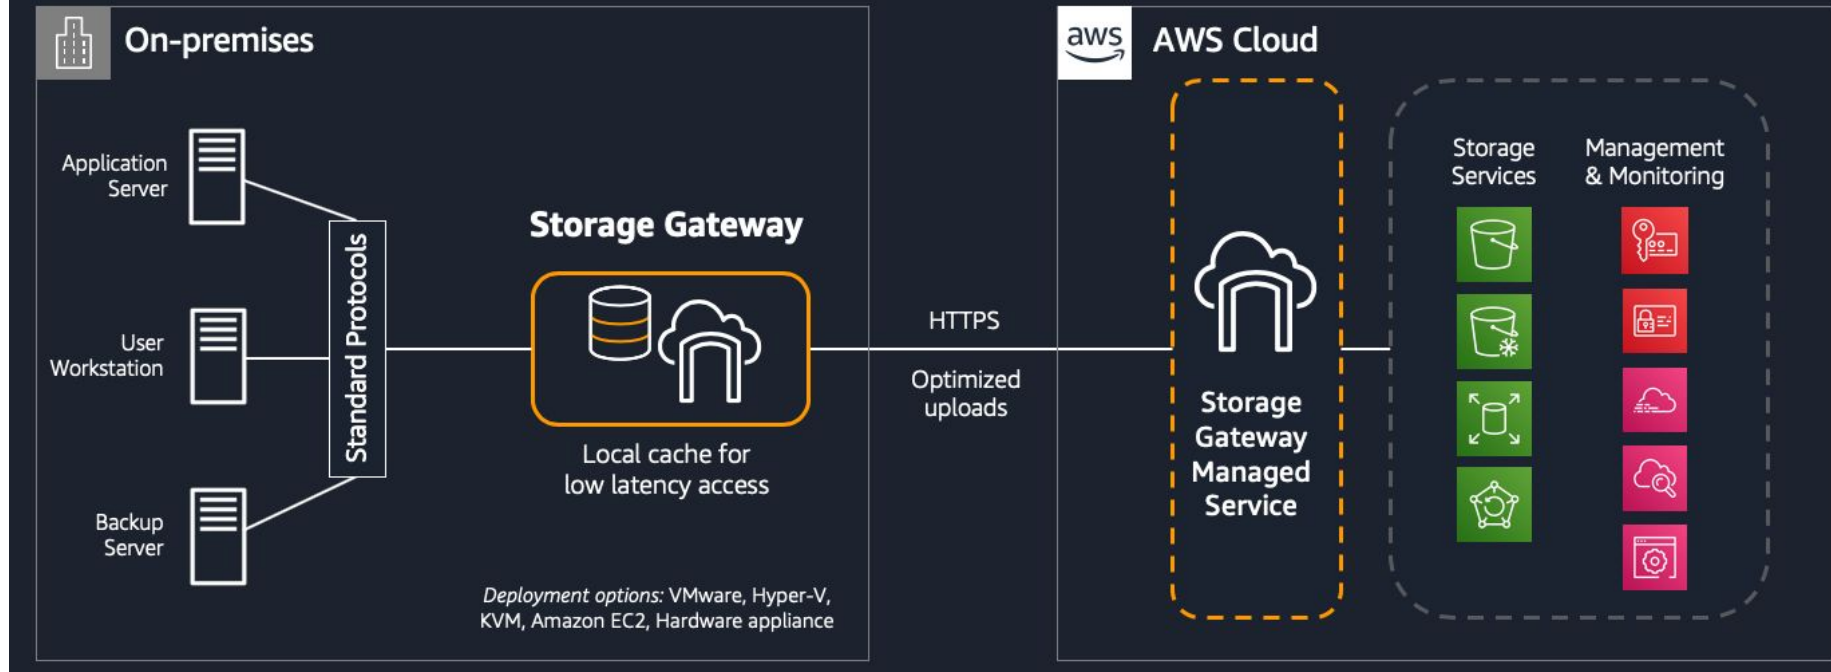
...but leverage the infinite scale of the cloud

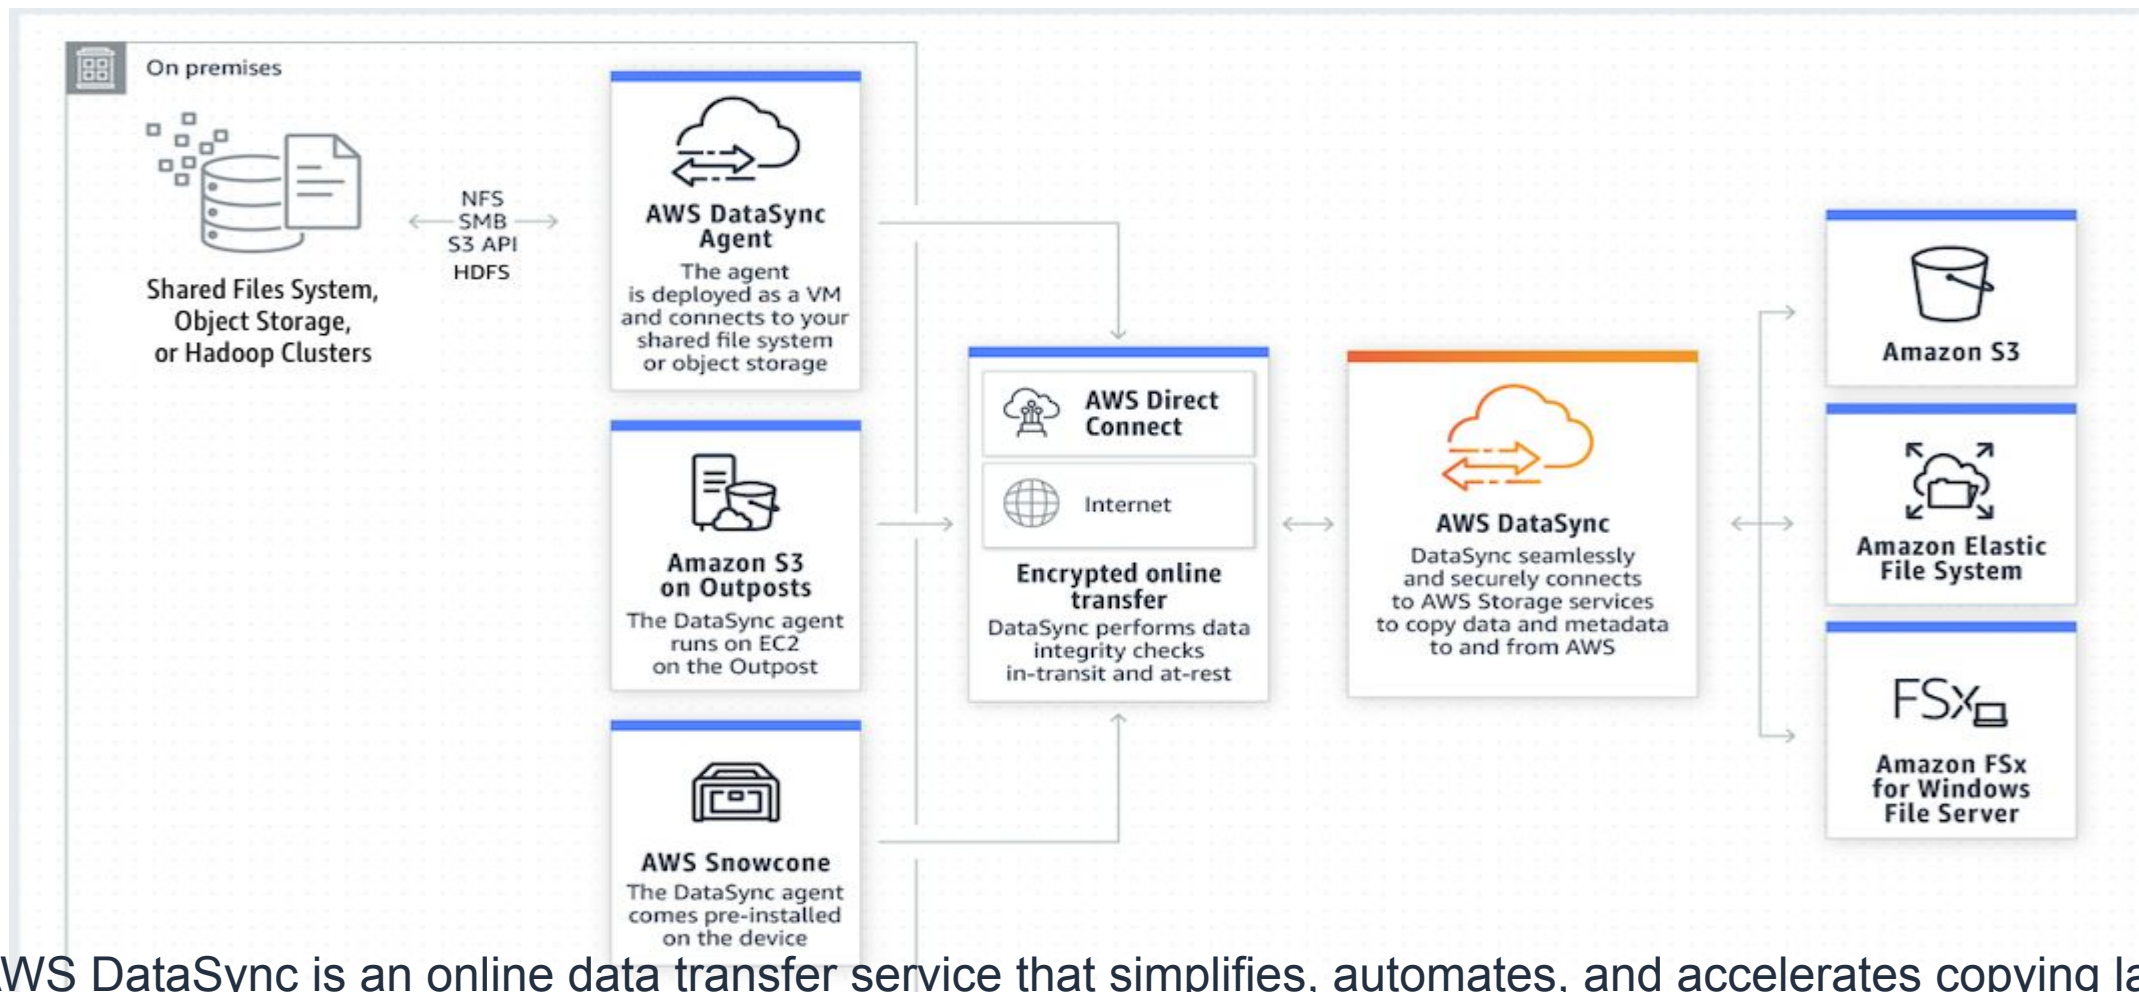
...while accessing the power of cloud compute and analytics

...but I want to manage & monitor from a single pane of glass

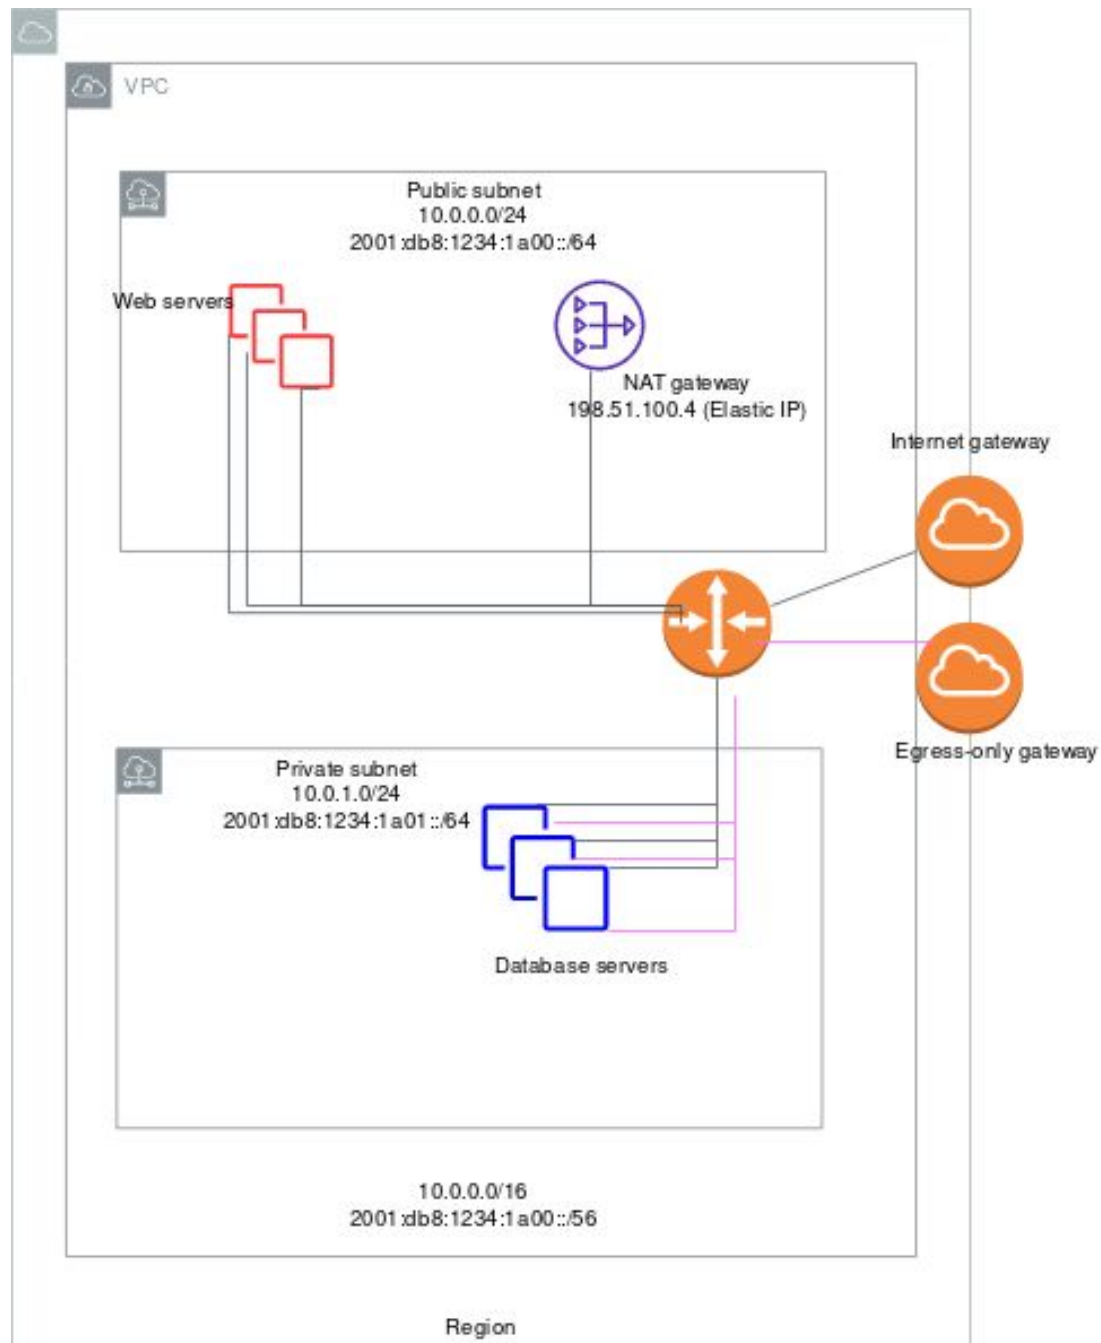
...even when my data is in many locations

# STORAGE GATEWAY





AWS DataSync is an online data transfer service that simplifies, automates, and accelerates copying large amounts of data between on-premises storage systems and AWS Storage services, as well as between AWS Storage services. DataSync can copy data between Network File System (NFS), Server Message Block (SMB) file servers, Hadoop Distributed File Systems (HDFS), self-managed object storage, [AWS Snowcone](#), Amazon Simple Storage Service ([Amazon S3](#)) buckets, Amazon Elastic File System ([Amazon EFS](#)) file systems, and [Amazon FSx for Windows File Server](#) file systems.



# CloudFront



- CloudFront is the **Content Delivery Network (CDN)** of AWS.
- Basically, it is a globally distributed network located on different geographical places.

# CloudFront



- The aim of this service is to **provide faster distribution** to the dynamic or the static web contents.
- World-wide distributed data centers which are called **edge locations** are the backbone of CloudFront.





## IAM roles

**You can use roles to delegate access to users, applications, or services that don't normally have access to your AWS resources.**

For example, you might want to grant users in your AWS account access to resources they don't usually have, or grant users in one AWS account access to resources in another account. Or you might want to allow a mobile app to use AWS resources, but not want to embed AWS keys within the app (where they can be difficult to rotate and where users can potentially extract them).

Sometimes you want to give AWS access to users who already have identities defined outside of AWS, such as in your corporate directory. Or, you might want to grant access to your account to third parties so that they can perform an audit on your resources.



# AWS Key Management Service

What is AWS Key Management Service?

AWS Key Management Service (AWS KMS) is a secured service to create and control the encryption keys. It is integrated with other AWS services such as Amazon EBS, Amazon S3 to provide data at rest security with encryption keys.

KMS is a global service but keys are regional which means you can't send keys

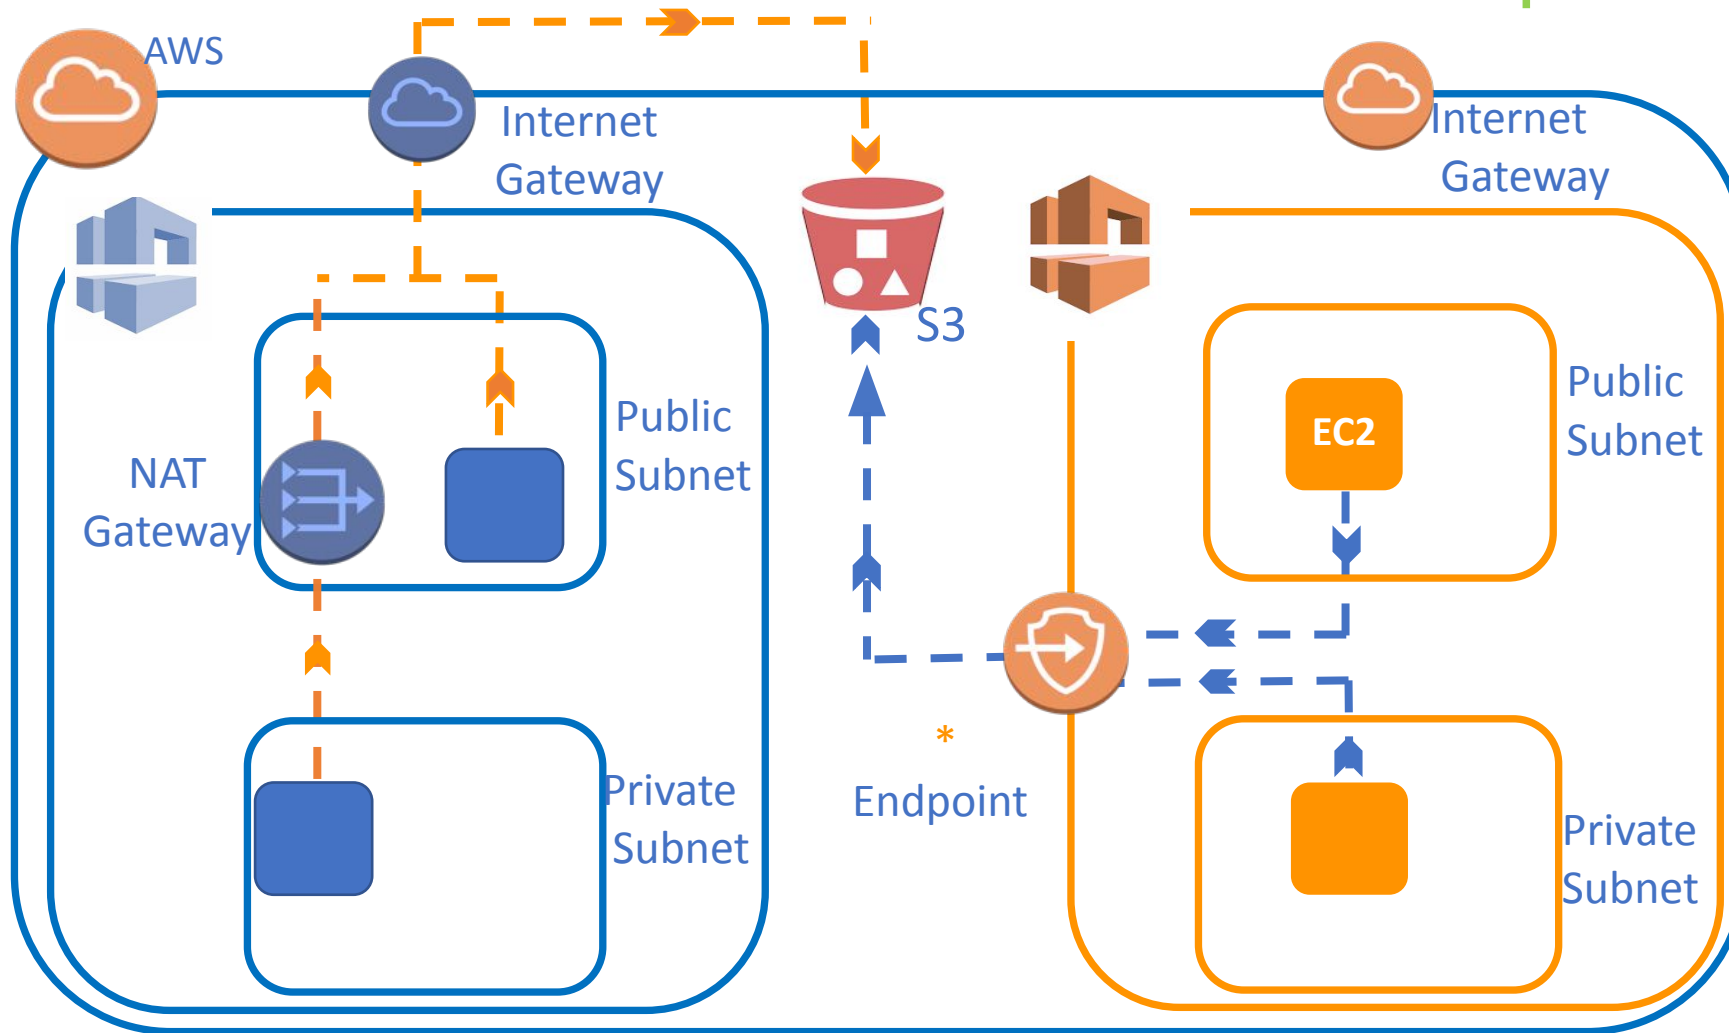
outside the region in which they are created.





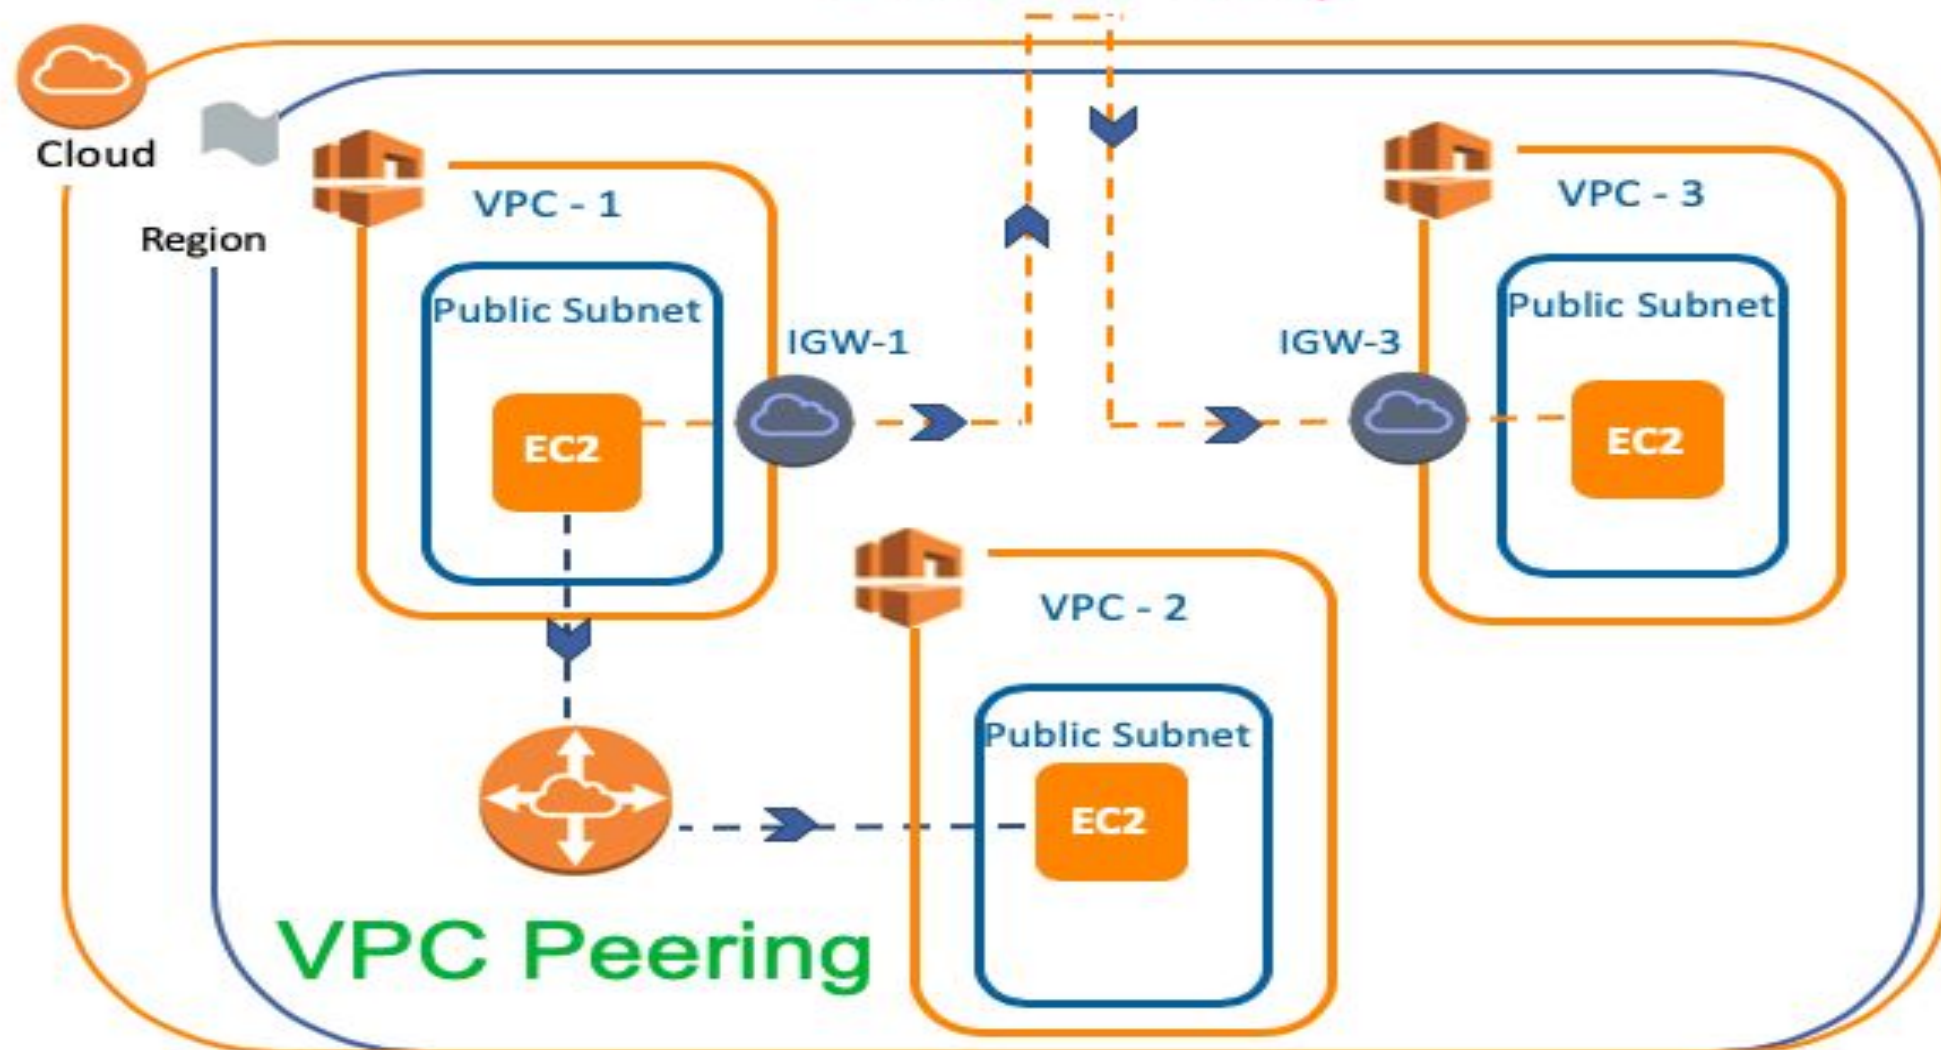
## Classic Way

## VPC Endpoint



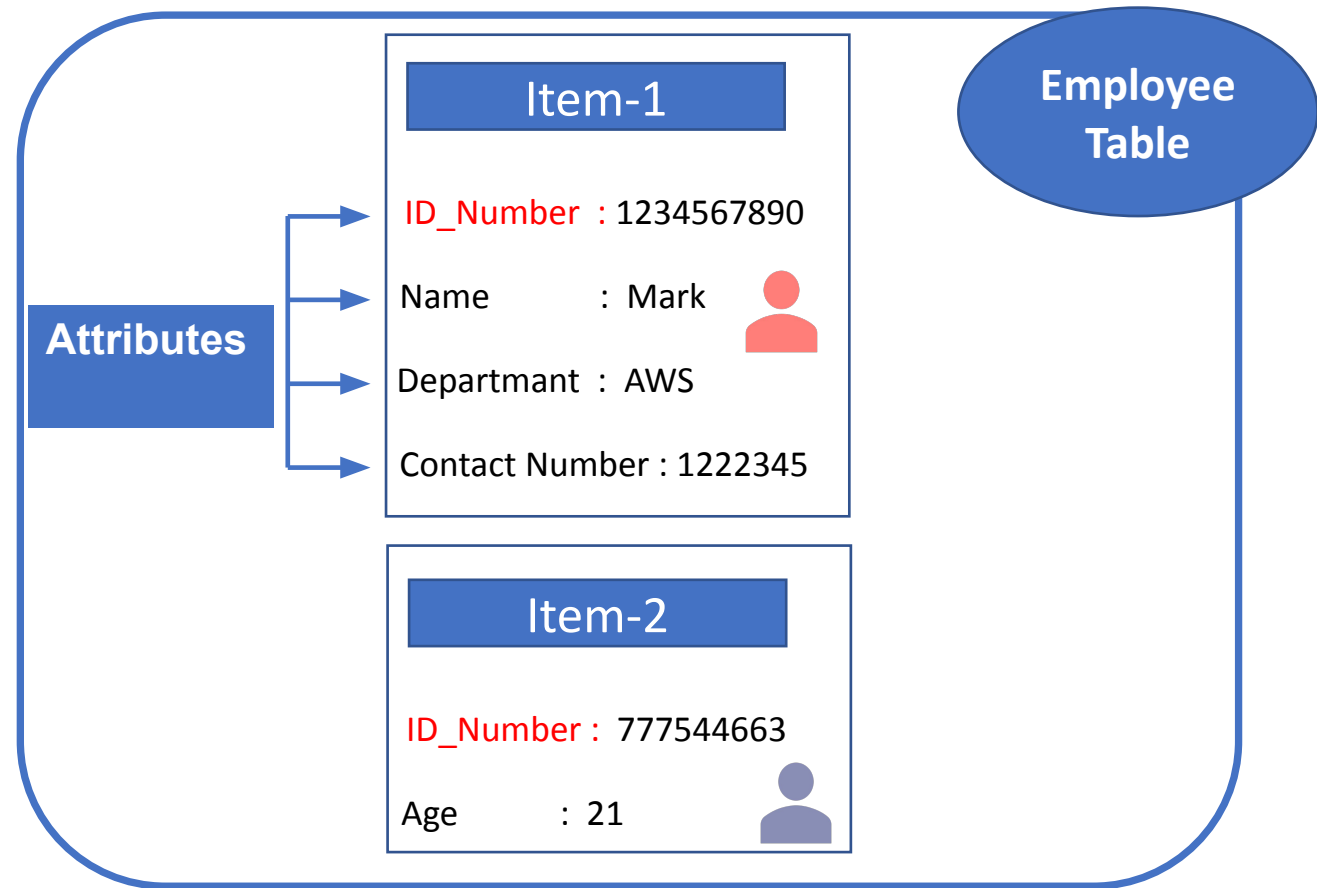


## Classic way





**amazon**  
DynamoDB



# Query String

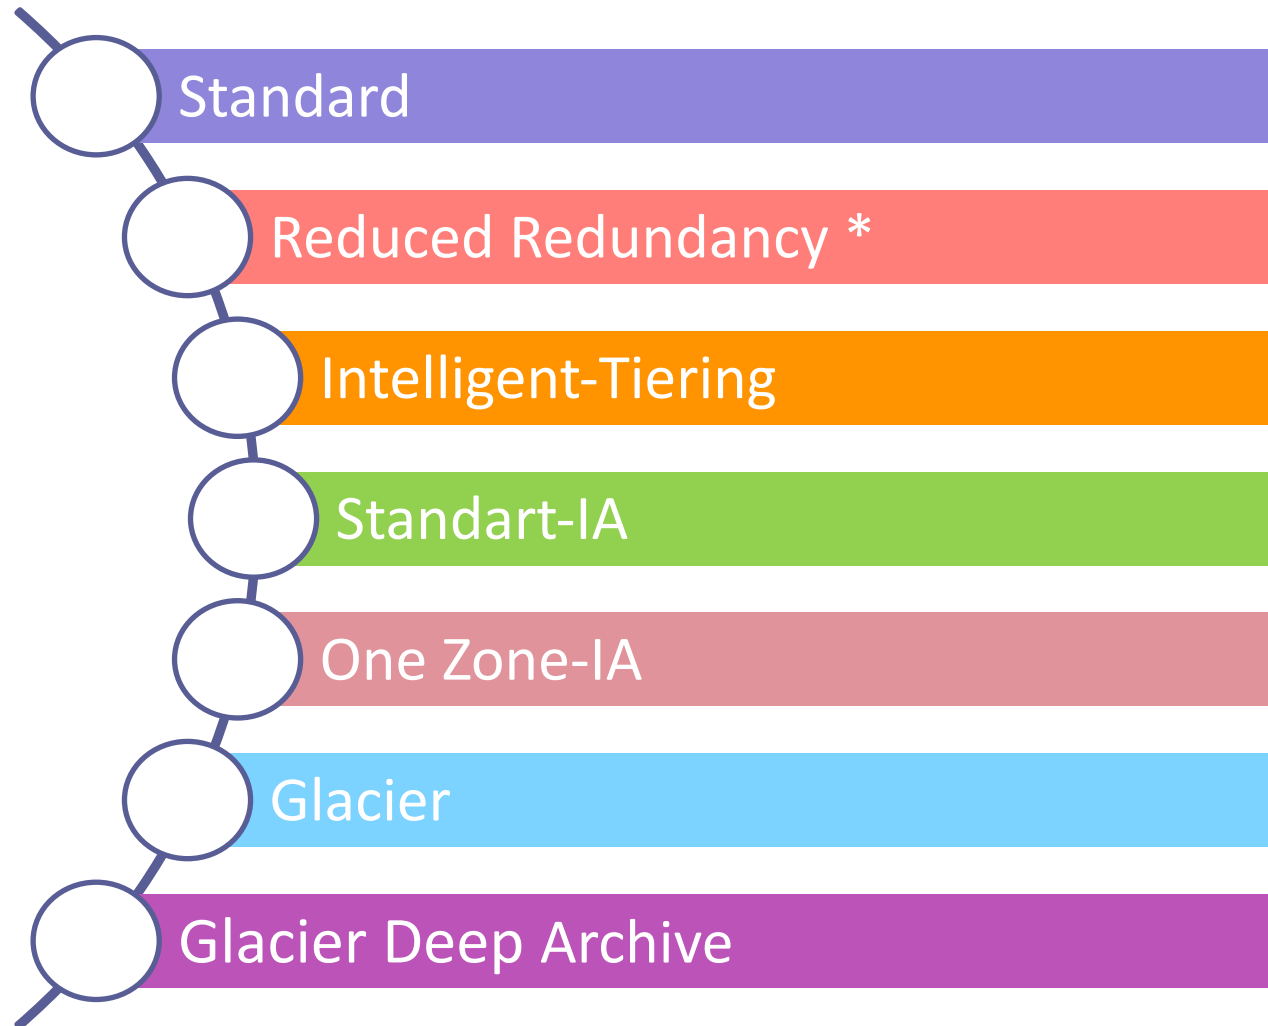
A **query string** is a part of a **uniform resource locator** (URL) that assigns values to specified parameters. A query string commonly includes fields added to a base URL by a Web browser or other client application, for example as part of an HTML, choosing the appearance of a page, or jumping to positions in multimedia content



Secure | [https://en.wikipedia.org/w/index.php?title=Query\\_string&action=edit](https://en.wikipedia.org/w/index.php?title=Query_string&action=edit)

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/QueryStringParameters.html>

# Storage Classes




\* Not recommended  
by AWS

# Storage Classes


## Intelligent Tier

Unpredictable  
Access Patterns




**Frequent  
Access Tier**

30 consecutive days after last access




**Infrequent  
Access Tier**

90 consecutive days after the last access



**Archive  
Access tier**

180 consecutive days after the last access

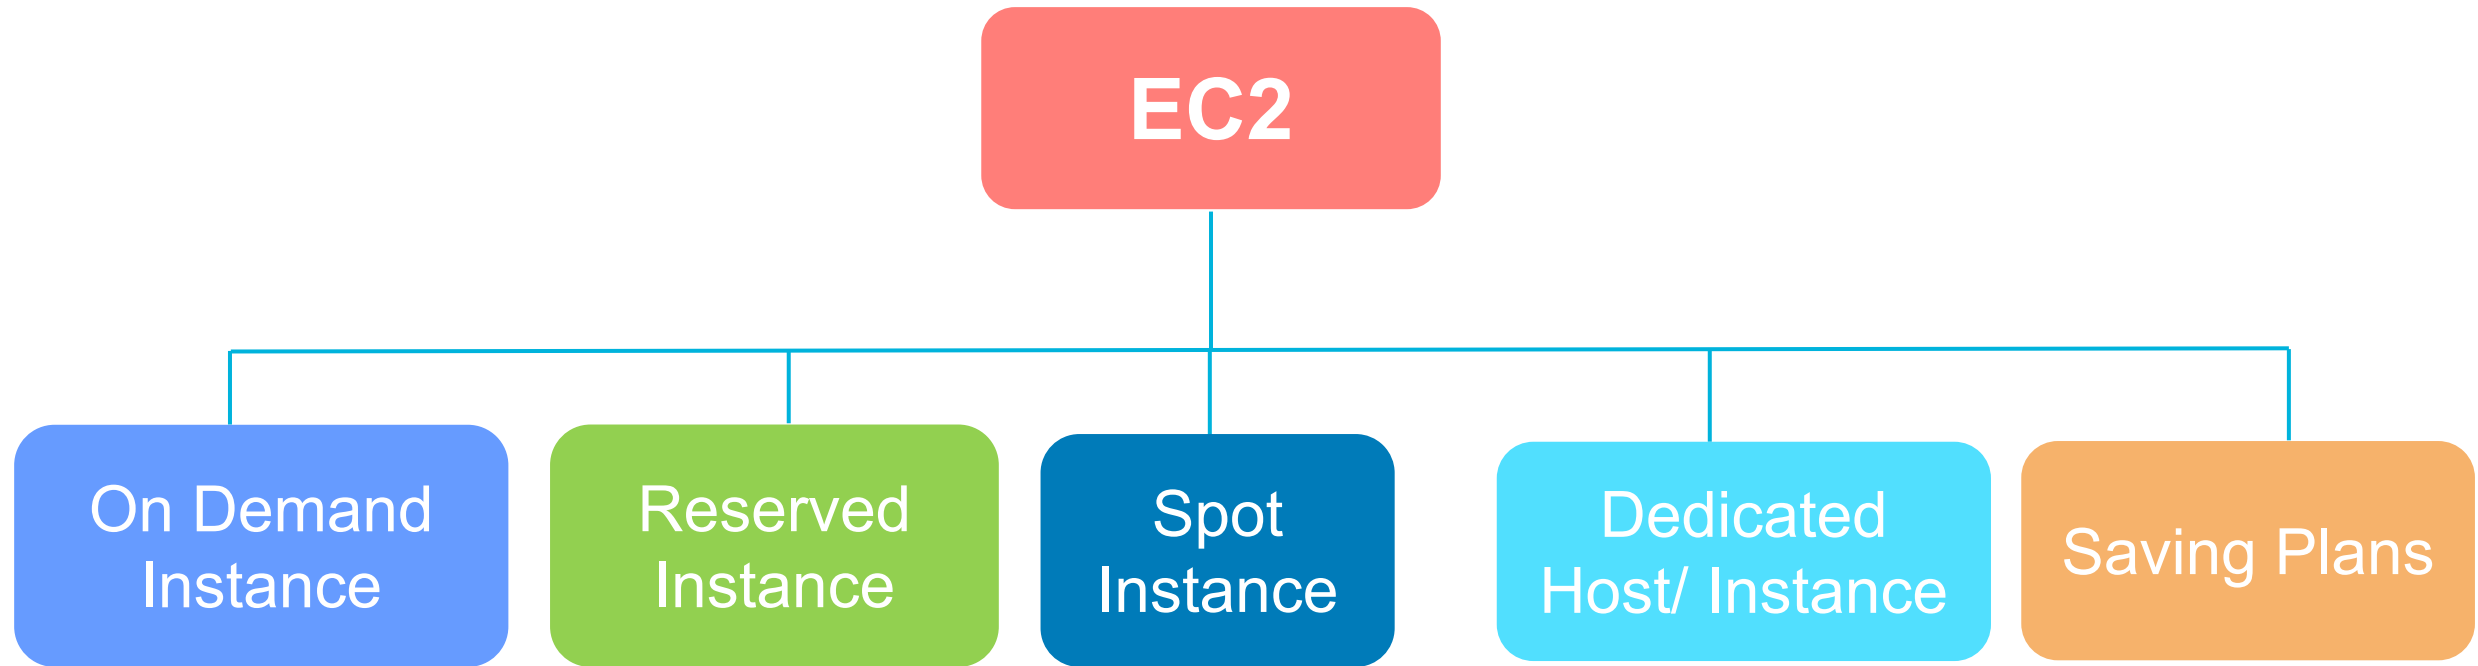


**Deep Archive  
Access tier**

- It is designed to optimize storage costs by automatically moving data to the most cost-effective storage access tier.
- There are 4 access tiers.
- It is ideal, if your access patterns are unknown or unpredictable.

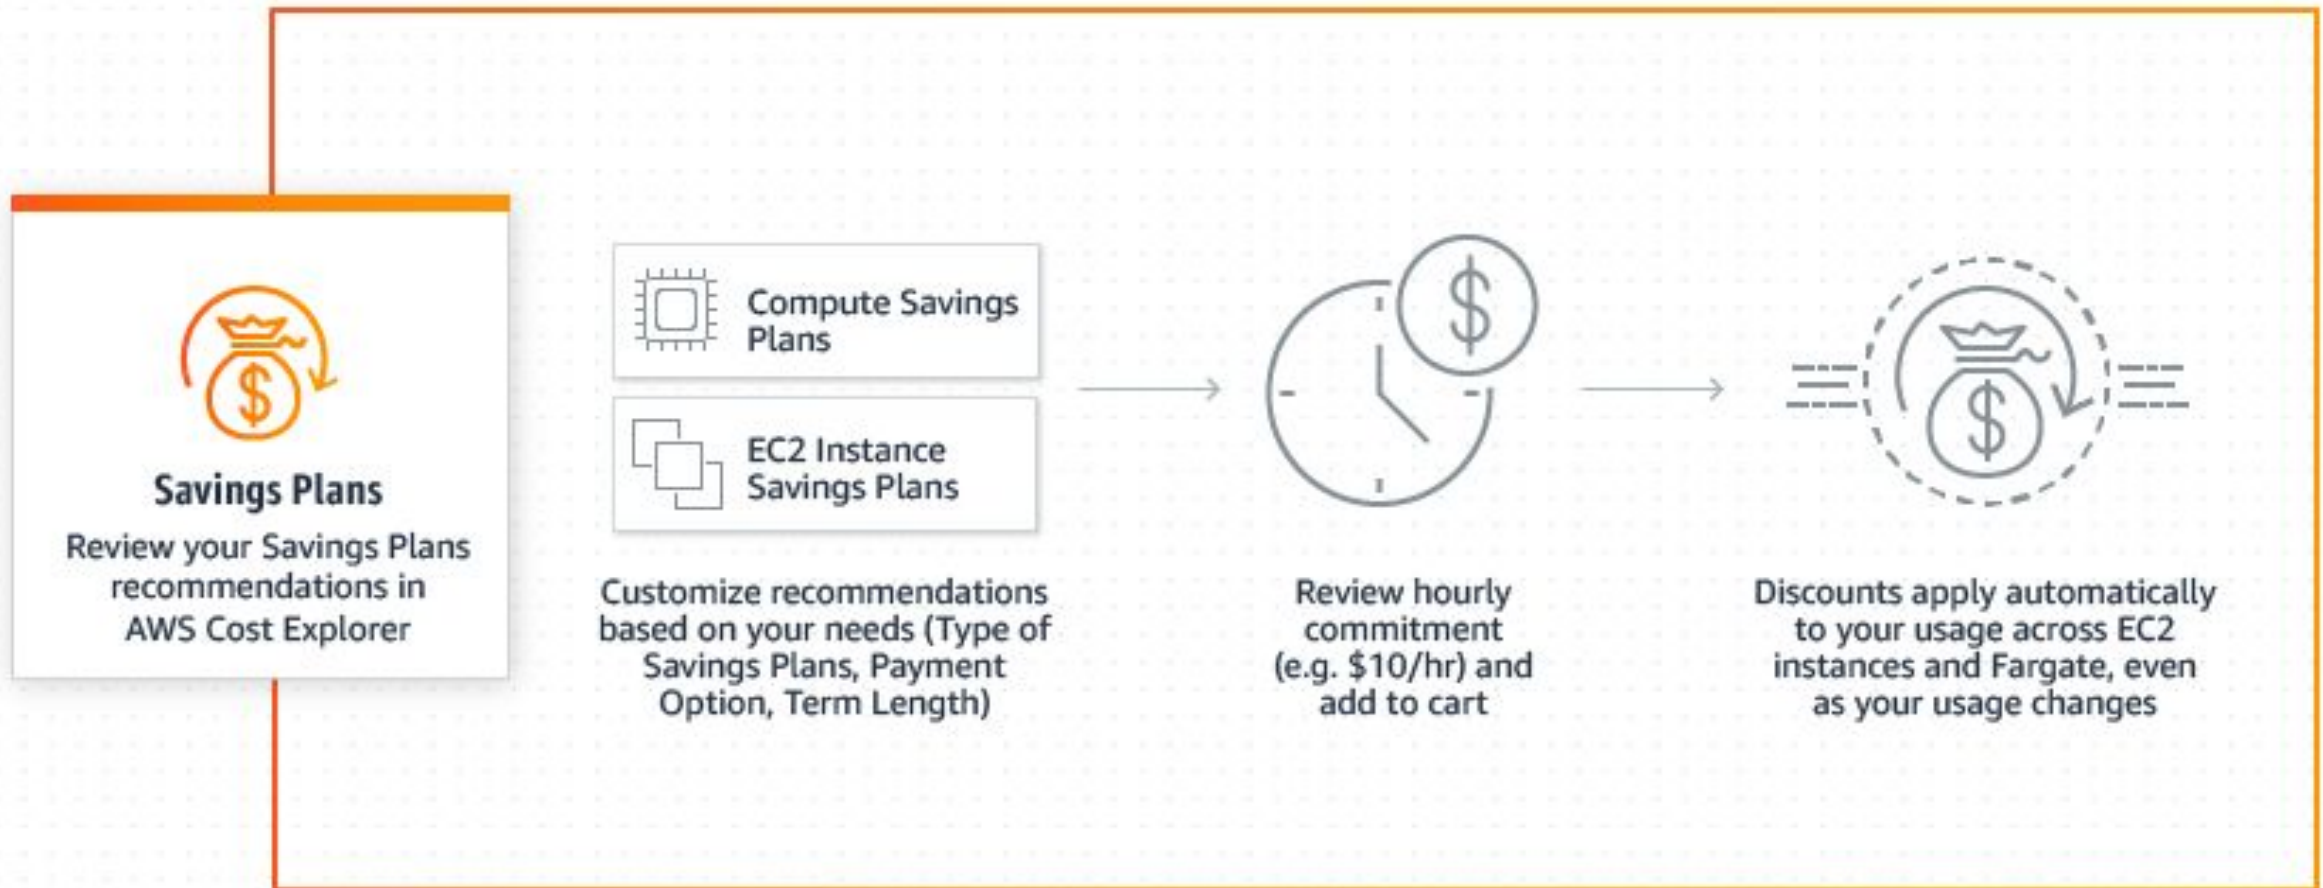


# Pricing Model of Instances






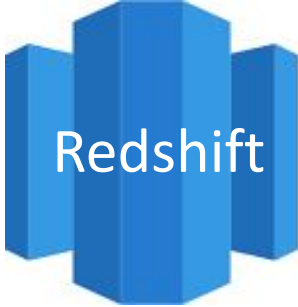
# Pricing Model of Instances






# Amazon Redshift



Database	Data Warehouse
 RDS	 Redshift
OLTP (Online Transaction Processing)	OLAP (Online Analytical Processing)
Used for data storing and managing	Used for data analyzing

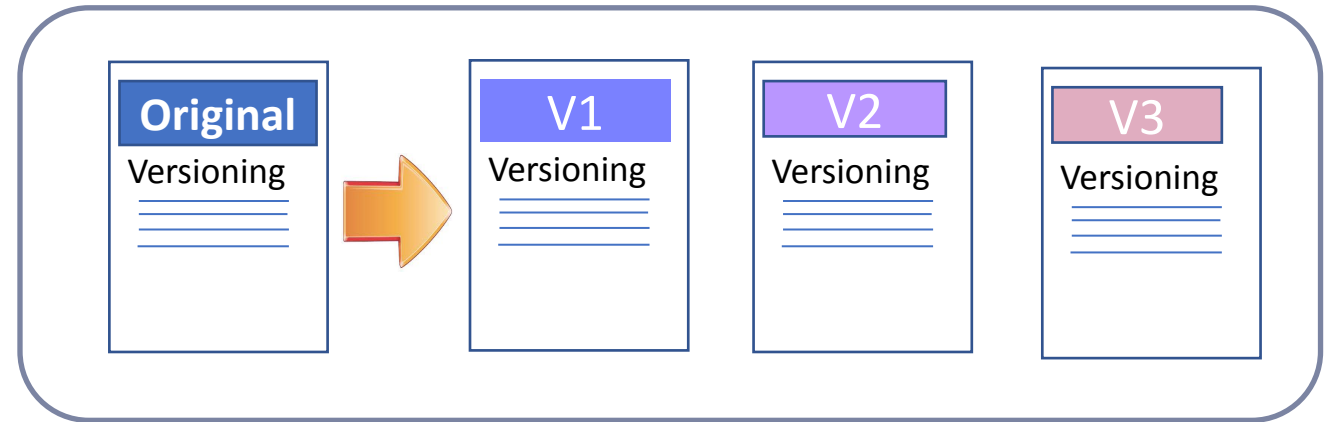
- Since the **analyzing process causes an extra workload on database** we prefer to use data warehouse
- Amazon Redshift is a fully managed, cloud-based, petabyte-scale **data warehouse** service by Amazon Web Services (AWS).
- Amazon Redshift is an efficient solution to **collect and store** all your data to **analyze**.

# AWS ElastiCache



- ElastiCache is an **In-Memory Cache service of AWS**.
- In-Memory Cache is a **temporary and fast** storage component. These components are used to reduce the workload of the main data storage device such as a database.
- AWS offers Redis and Memcached in-memory cache option which are popular in market.

# Versioning



- Versioning is a way to **keep multiple versions** (deleted and changed versions) of an object in a bucket.
- By using versioning, all **unwanted user behavior and program errors** can be quickly recovered.



	General Purpose SSD		Provisioned IOPS SSD		
Volume type	gp3	gp2	io2 Block Express ‡	io2	io1
Durability	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)	99.999% durability (0.001% annual failure rate)	99.999% durability (0.001% annual failure rate)	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)
Use cases	<ul style="list-style-type: none"><li>• Low-latency interactive apps</li><li>• Development and test environments</li></ul>		Workloads that require: <ul style="list-style-type: none"><li>• Sub-millisecond latency</li><li>• Sustained IOPS performance</li><li>• More than 64,000 IOPS or 1,000 MiB/s of throughput</li></ul>	<ul style="list-style-type: none"><li>• Workloads that require sustained IOPS performance or more than 16,000 IOPS</li><li>• I/O-intensive database workloads</li></ul>	




# Placement groups


## Cluster placement groups



## spread placement group



## Partition placement groups





## What is Amazon FSx for Windows File Server?

- Amazon FSx for Windows File Server is an FSx solution that offers a scalable and shared file storage system on the Microsoft Windows server.
- Using the Server Message Block (SMB) protocol with Amazon FSx Can access file storage systems from multiple windows servers.
- It offers to choose from HDD and SSD storage, offers high throughput, and IOPS with sub-millisecond latencies for Windows workloads.
- Using SMB protocol, Amazon FSx can connect file systems to Amazon EC2, Amazon ECS, Amazon WorkSpaces, Amazon AppStream 2.0 instances, and on-premises servers using AWS Direct Connect or AWS VPN.
- It provides high availability (Multi-AZ deployments) with an active and standby file server in separate AZs.
- It automatically and synchronously replicates data in the standby Availability Zone (AZ) to manage failover.
- Using AWS DataSync with Amazon FSx helps to migrate self-managed file systems to Windows storage systems.
- It offers identity-based authentication using Microsoft Active Directory (AD).
- It automatically encrypts data at rest with the help of AWS Key Management Service (AWS KMS). It uses SMB Kerberos session keys to encrypt data in transit.





## What is Amazon FSx for Lustre?

- Amazon FSx for Lustre is an FSx solution that offers scalable storage for the Lustre system (parallel and high-performance file storage system).
- It supports fast processing workloads like custom electronic design automation (EDA) and high-performance computing (HPC).
- It provides shared file storage with hundreds of gigabytes of throughput, sub-millisecond latencies, and millions of IOPS.
- It offers a choice between SSD and HDD for storage.
- It integrates with Amazon S3 to process data concurrently using parallel data-transfer techniques.
- It stores datasets in S3 as files instead of objects and automatically updates with the latest data to run the workload.
- It offers to select unreplicated file systems for shorter-term data processing.
- It can be used with existing Linux-based applications without any changes.
- It offers network access control using POSIX permissions or Amazon VPC Security Groups.
- It easily provides data-at-rest and in-transit encryption.
- AWS Backup can also be used to backup Lustre file systems.
- It integrates with SageMaker to process machine learning workloads.

# AWS Lambda

## Lambda and API Gateway

- Alexa
- **API Gateway**
- CloudTrail
- CloudWatch Events
- CloudWatch Logs
- CloudFormation
- CloudFront (Lambda@Edge)
- CodeCommit
- CodePipeline
- Cognito
- Config
- DynamoDB
- EC2
- ElastiCache
- Elastic Load Balancing
- EFS
- IoT
- IoT Events
- Kinesis Firehose
- Kinesis Streams
- Lex
- RDS
- S3
- S3 Batch
- SES
- SNS
- SQS
- Step Functions
- X-Ray



Although you can trigger Lambda function in different way, the most common trigger of the Lambda function is **API Gateway**

Thanks to Lambda and API Gateway combination you can create fully automated system.

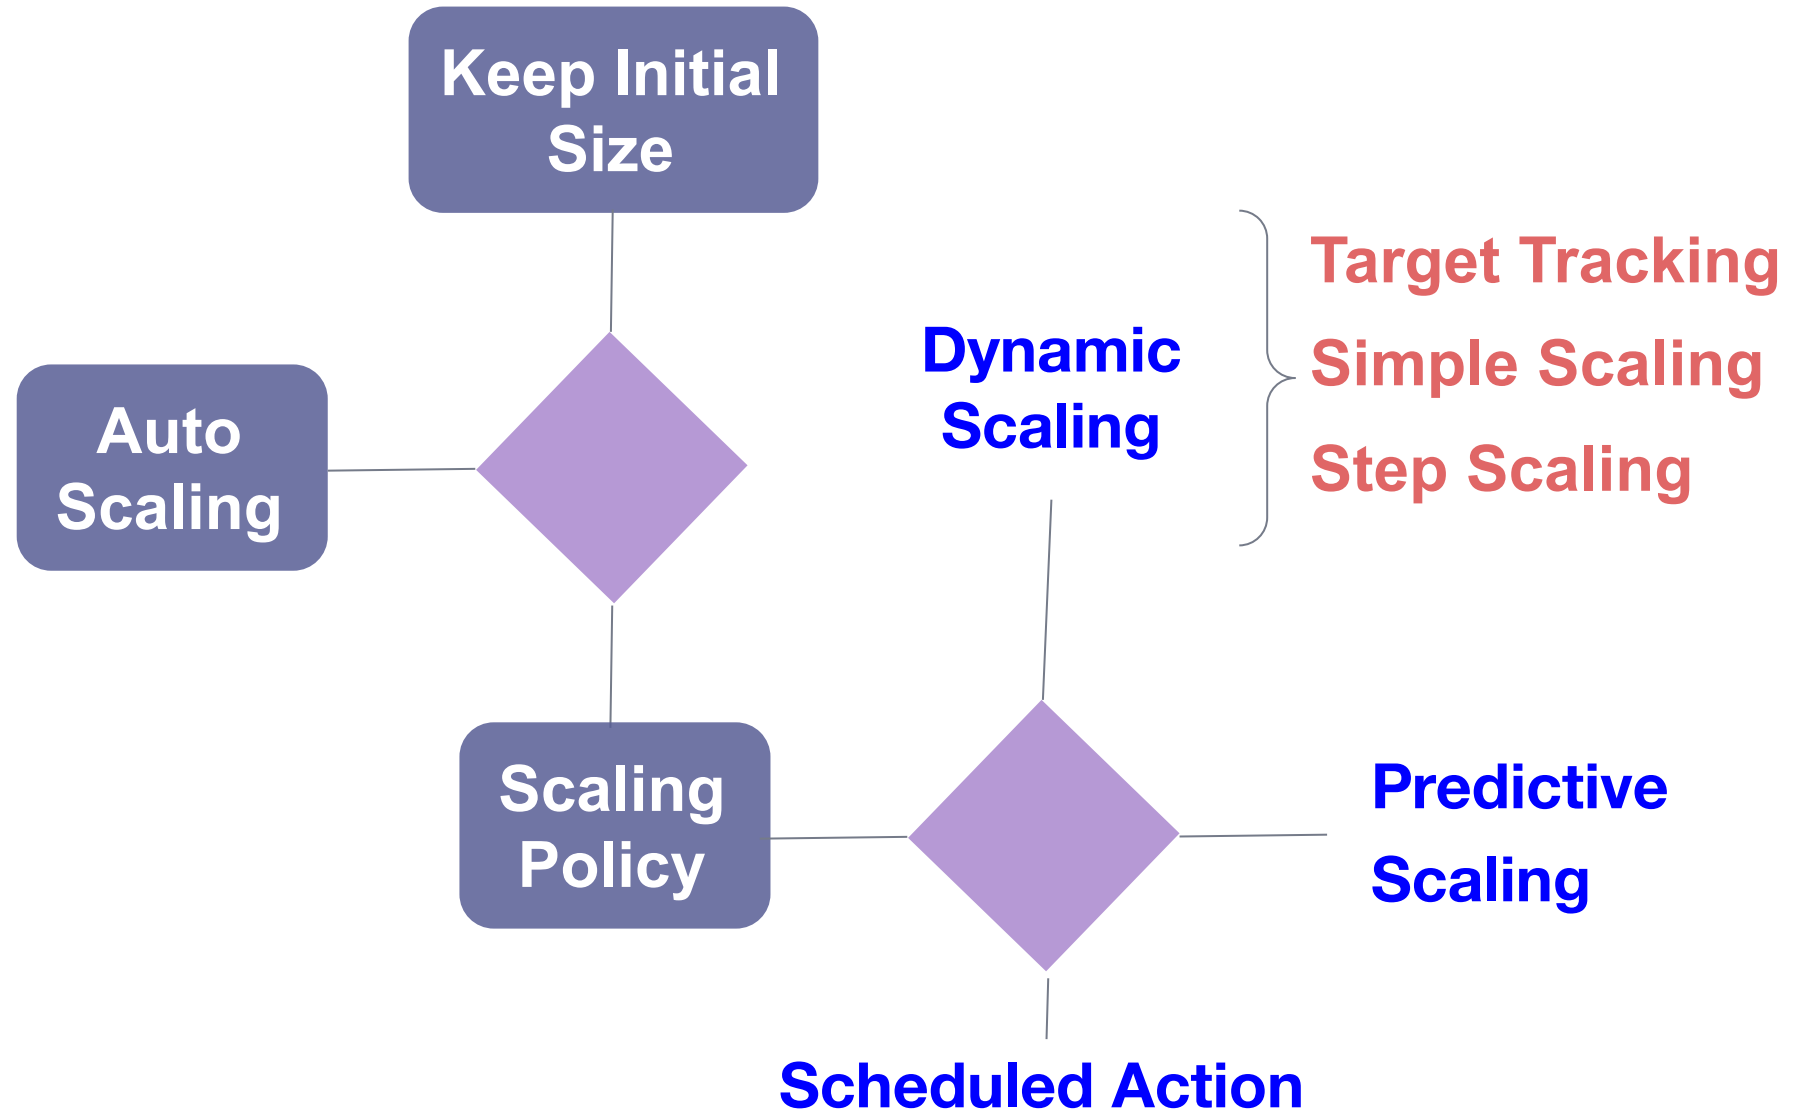






# Auto Scaling

## Auto Scaling Options





## What is AWS CloudTrail?

AWS CloudTrail is defined as a global service that permits users to enable operational and risk auditing of the AWS account.

It allows users to view, search, download, archive, analyze, and respond to account activity across the AWS infrastructure.

It records actions as an event taken by a user, role, or an AWS service in the AWS Management Console, AWS Command Line Interface, and AWS SDKs and APIs.

### **AWS CloudTrail mainly integrates with:**

- Amazon S3 can be used to retrieve log files.
- Amazon SNS can be used to notify about log file delivery to the bucket with Amazon Simple Queue Service (SQS).
- Amazon CloudWatch for monitoring and AWS Identity and Access Management (IAM) for security.



CloudTrail events of the past 90 days recorded by CloudTrail can be viewed in the CloudTrail console and can be downloaded in CSV or JSON file.

Trail log files can be aggregated from multiple accounts to a single bucket and can be shared between accounts.

AWS CloudTrail Insights enables AWS users to identify and respond to unusual activities of API calls by analyzing CloudTrail management events.

**There are three types of CloudTrail events:**

- **Management events or control plane operations**

- Example - Amazon EC2 CreateSubnet API operations and CreateDefaultVpc API operations


- **Data events**

- Example - S3 Bucket GetObject, DeleteObject, and PutObject API operations

- **CloudTrail Insights events (unusual activity events)**

- Example - Amazon S3 deleteBucket API, Amazon EC2 AuthorizeSecurityGroupIngress API

# AWS ElastiCache



- ElastiCache is an **In-Memory Cache service of AWS**.
- In-Memory Cache is a **temporary and fast** storage component. These components are used to reduce the workload of the main data storage device such as a database.
- AWS offers Redis and Memcached in-memory cache option which are popular in market.


# AWS ElastiCache

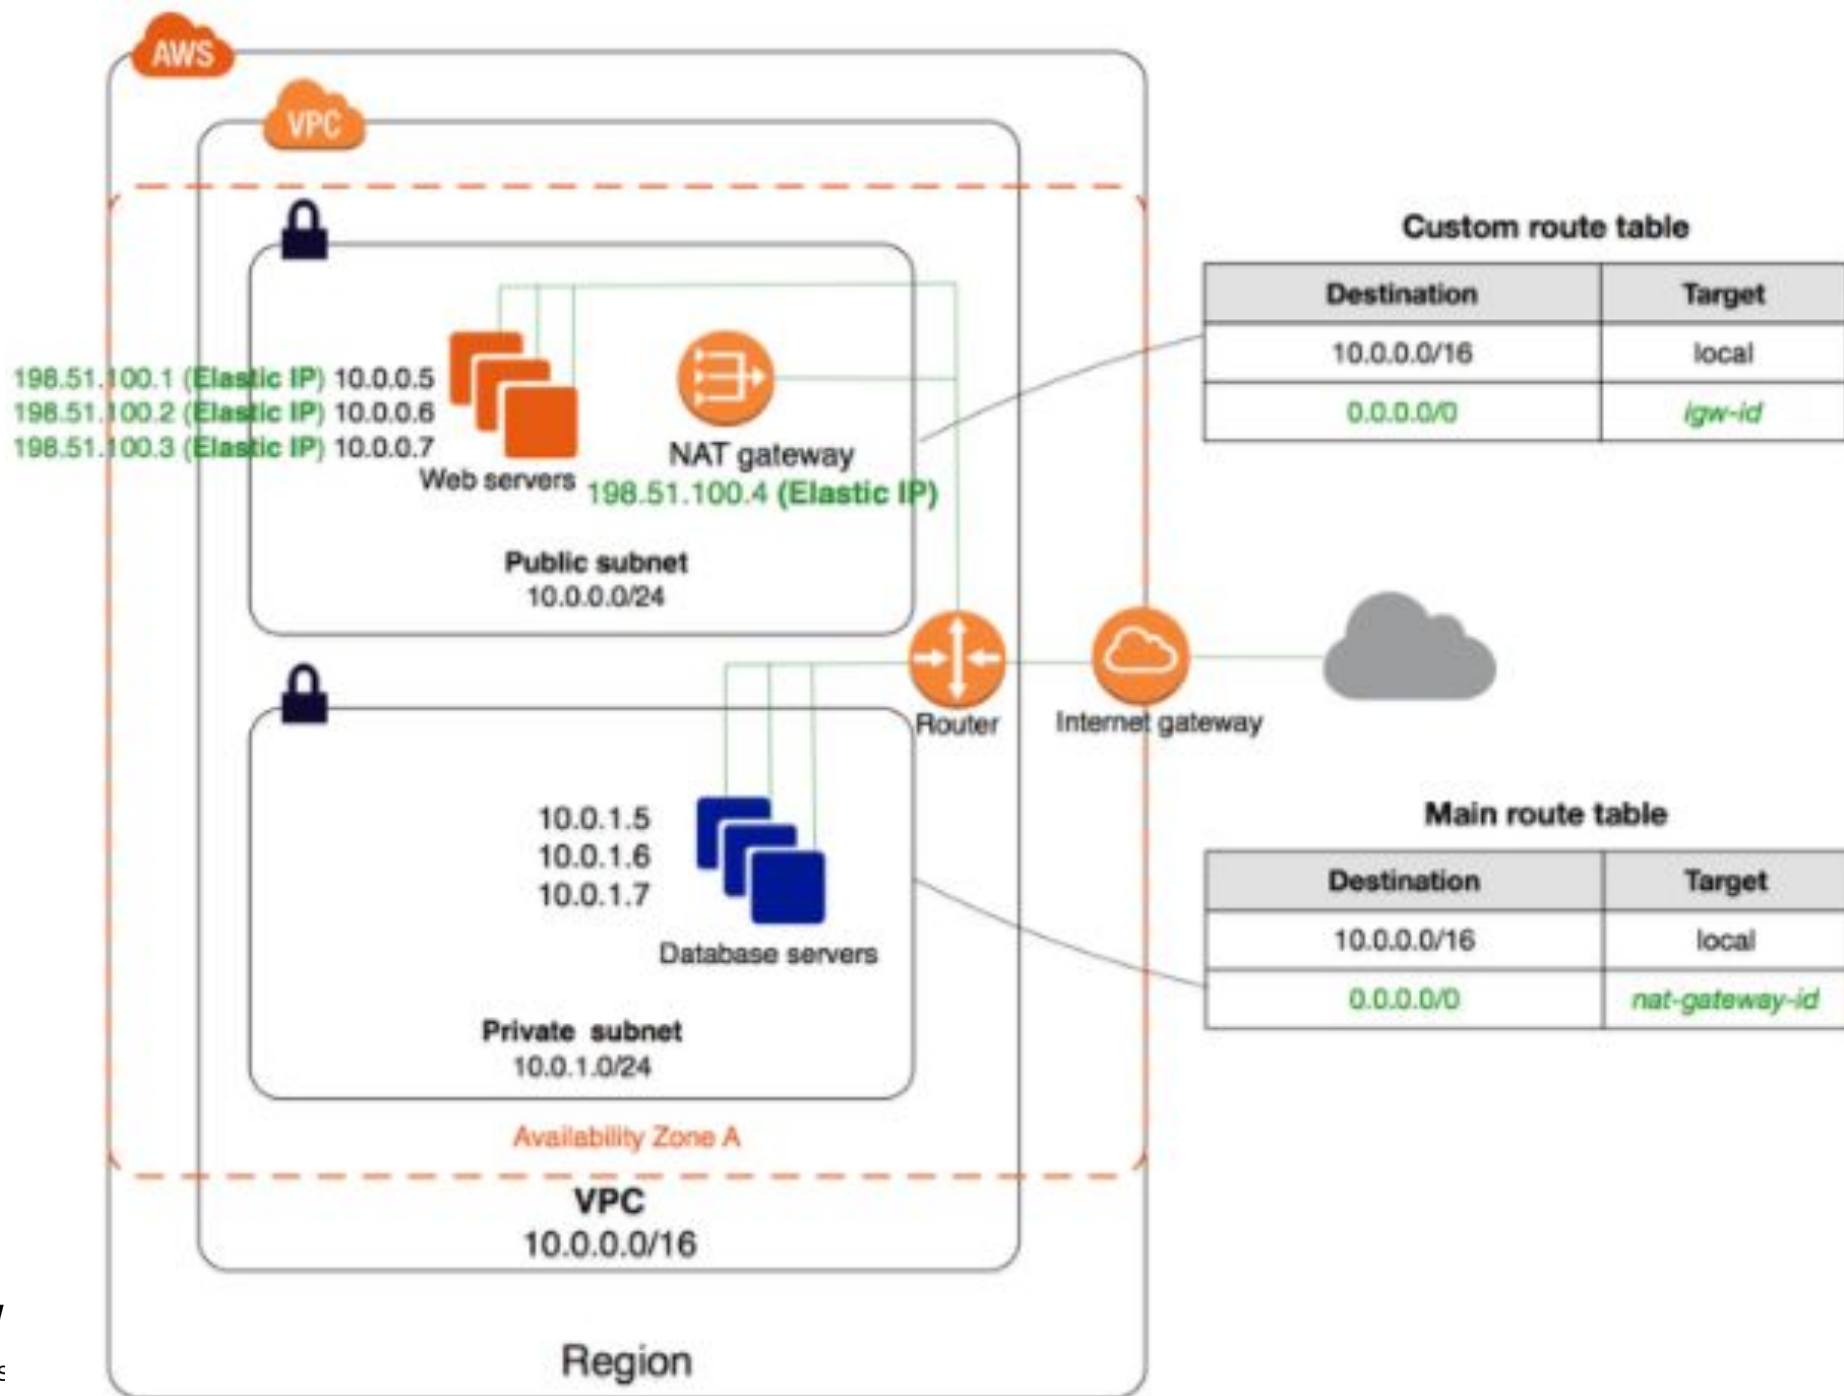


After ElastiCache - **First Query**

After ElastiCache - **Second Query**

Before ElastiCache






# Routing Policies of Amazon Route 53



## Routing Policies



- Routing Policy is a component that is used for configuring the traffic **according to different types of scenarios.**
- How traffic is routed depends on which **topic the policy is based on.**



# Routing Policies of Amazon Route 53

## Simple Routing Policy

**NO HEALTH CHECK**



**APPLICATION**

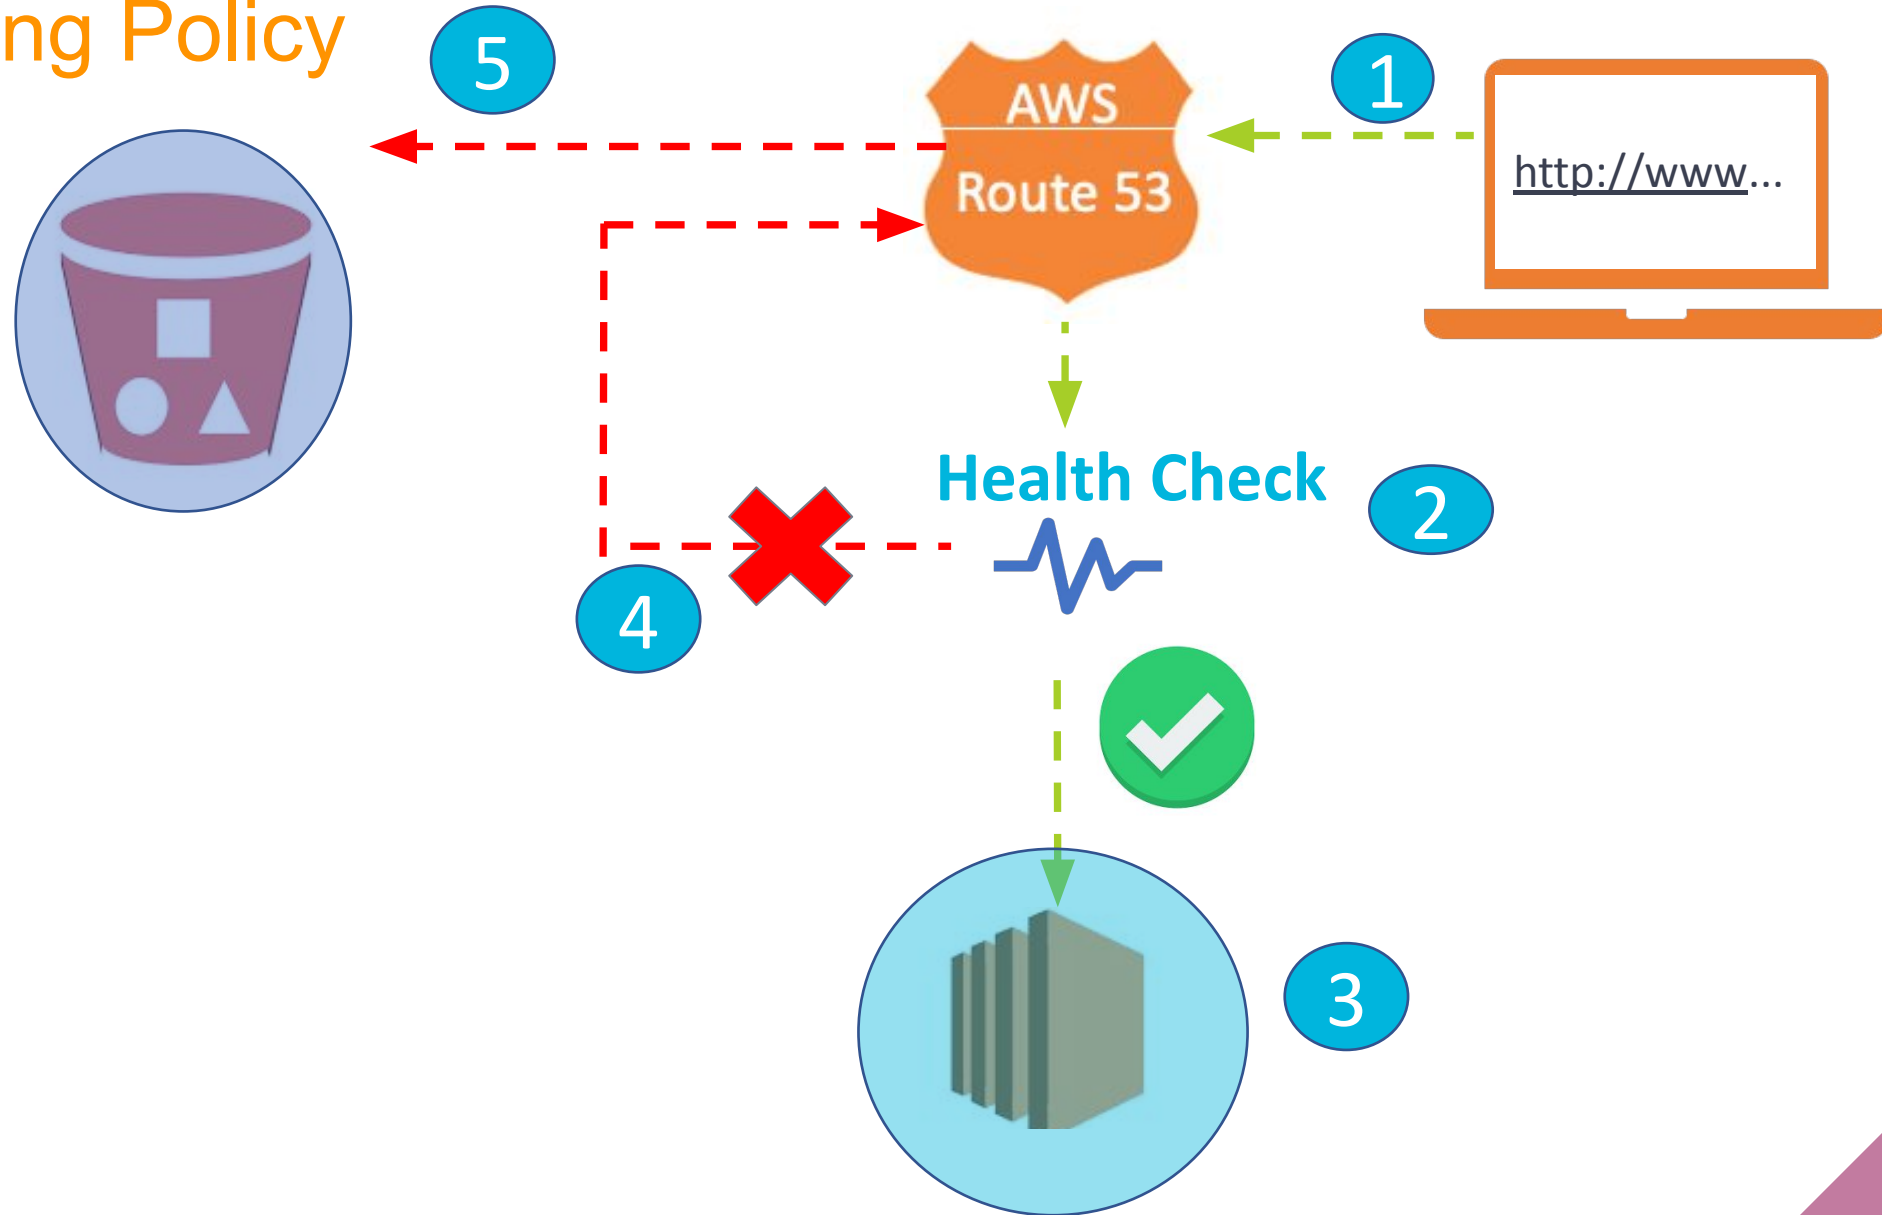
A: 11.22.33.44





# Routing Policies of Amazon Route 53

## Failover Routing Policy





# Routing Policies of Amazon Route 53


## Weighted Routing Policy





# Routing Policies of Amazon Route 53


## Latency Routing Policy





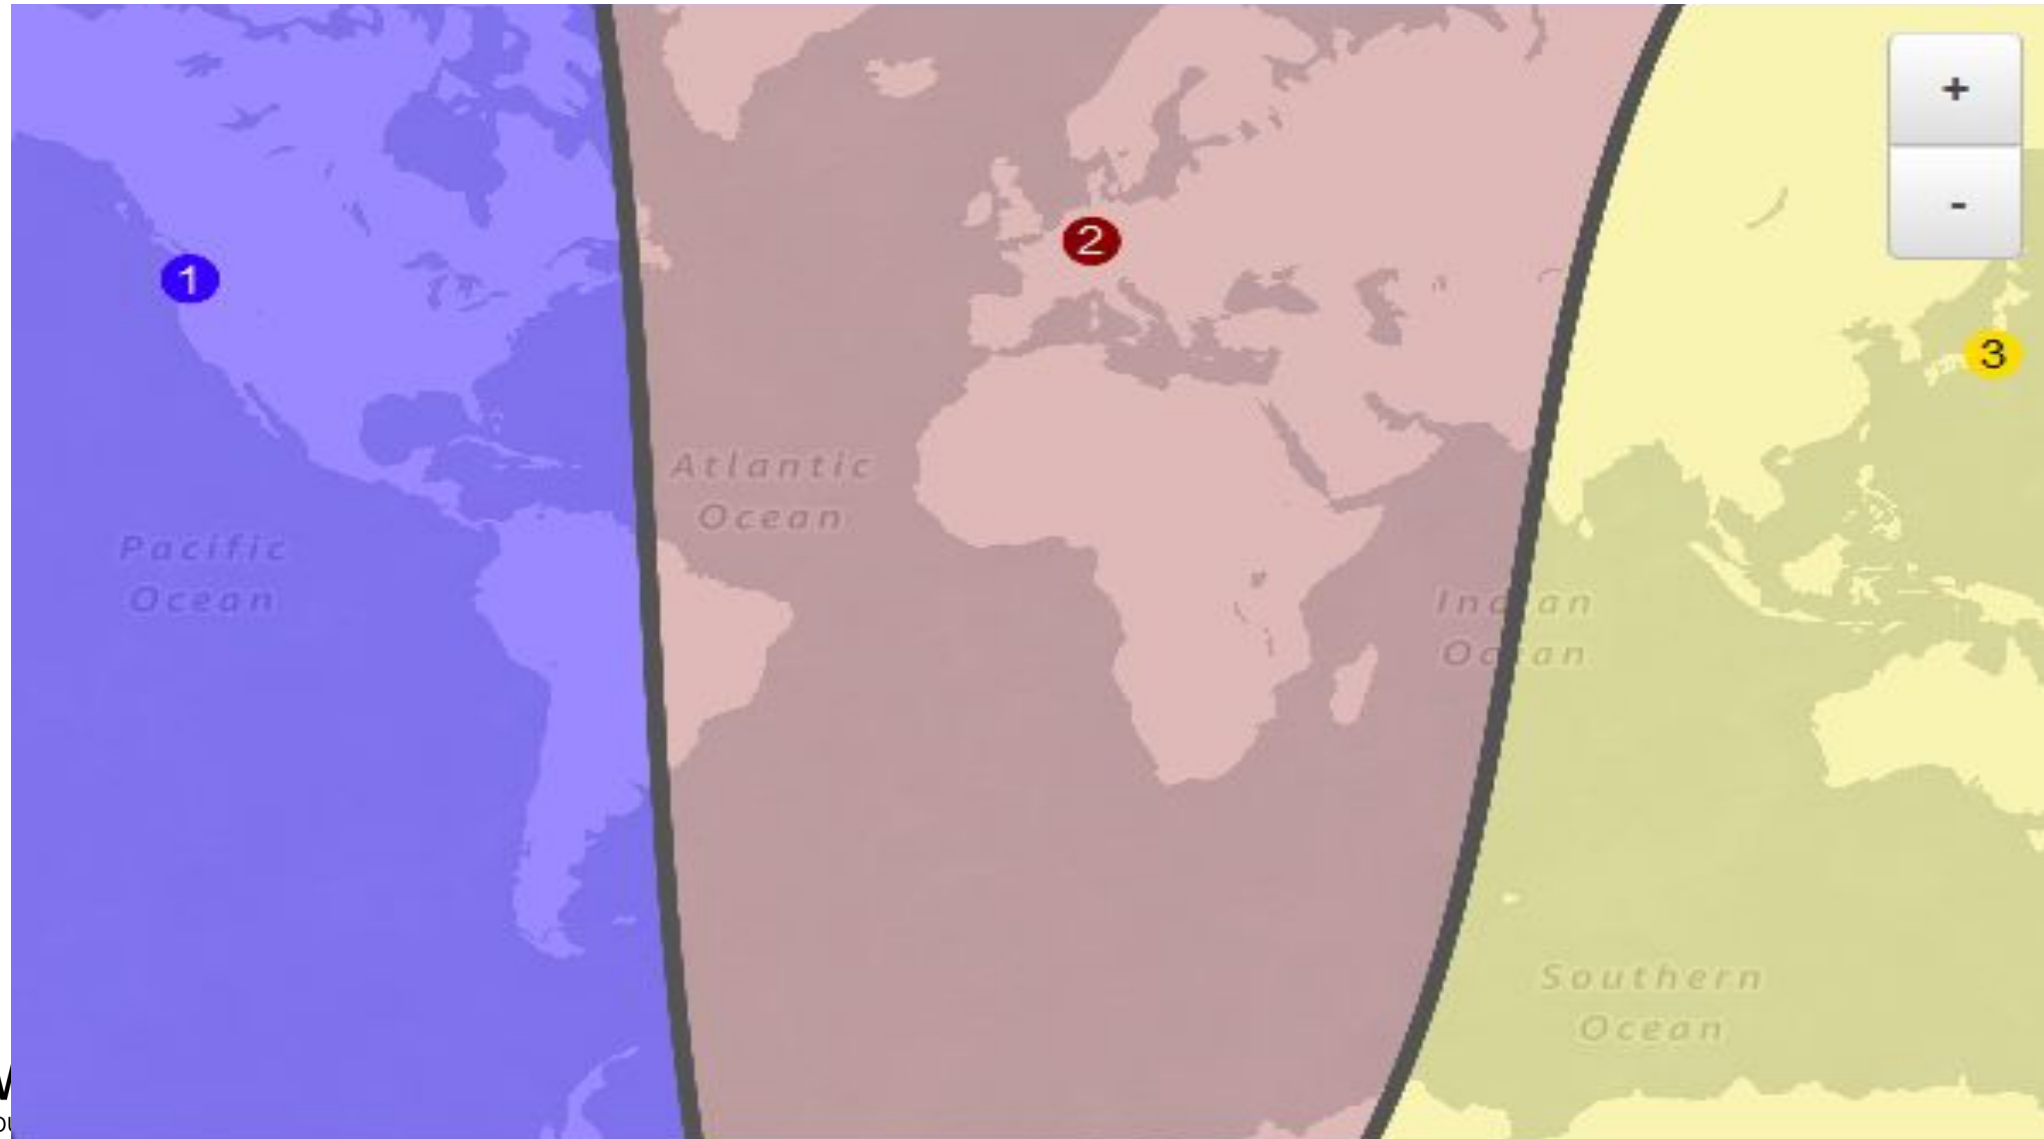
# Routing Policies of Amazon Route 53

## Geolocation Routing Policy



# Routing Policies of Amazon Route 53

## Geo Proximity Routing Policy





# Routing Policies of Amazon Route 53

## Multivalue Answer Routing Policy





**AZ-1**

**AZ-2**

**AZ-3**

<b>A</b>	<b>24</b>	<b>24</b>	<b>NO CHOICE</b>
<b>B</b>	<b>12</b>	<b>12</b>	<b>12</b>
<b>C</b>	<b>12</b>	<b>12</b>	<b>NO CHOICE</b>
<b>D</b>	<b>8</b>	<b>8</b>	<b>8</b>





# THANKS!

## Any questions?

You can find me at:

- ▶ @armando
- ▶ armando@clarusway.com

